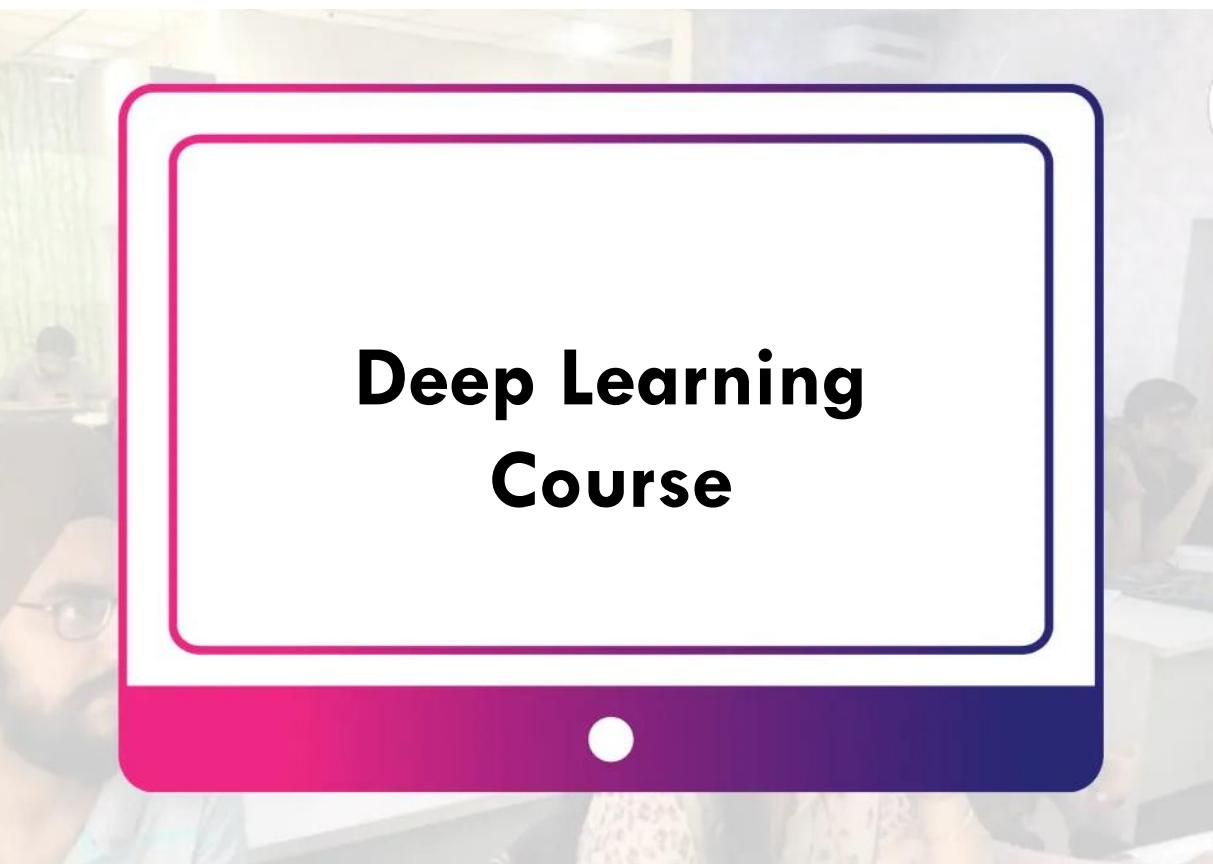
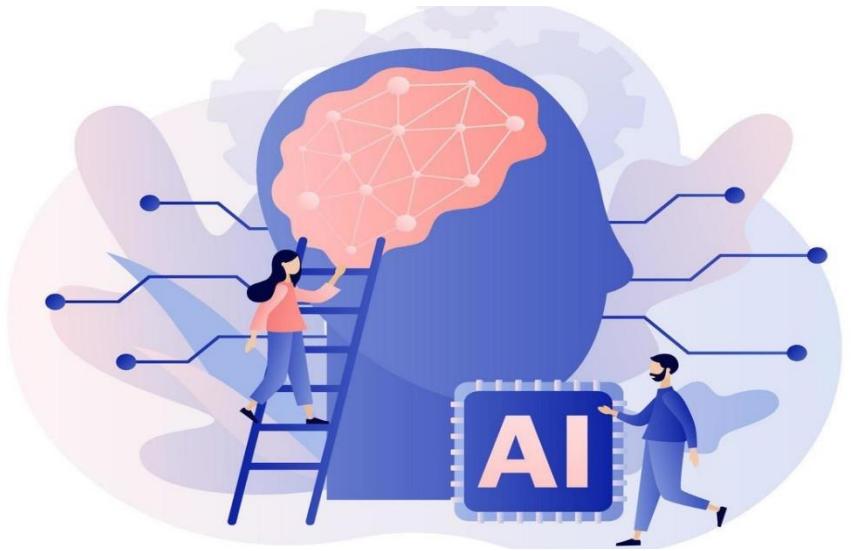




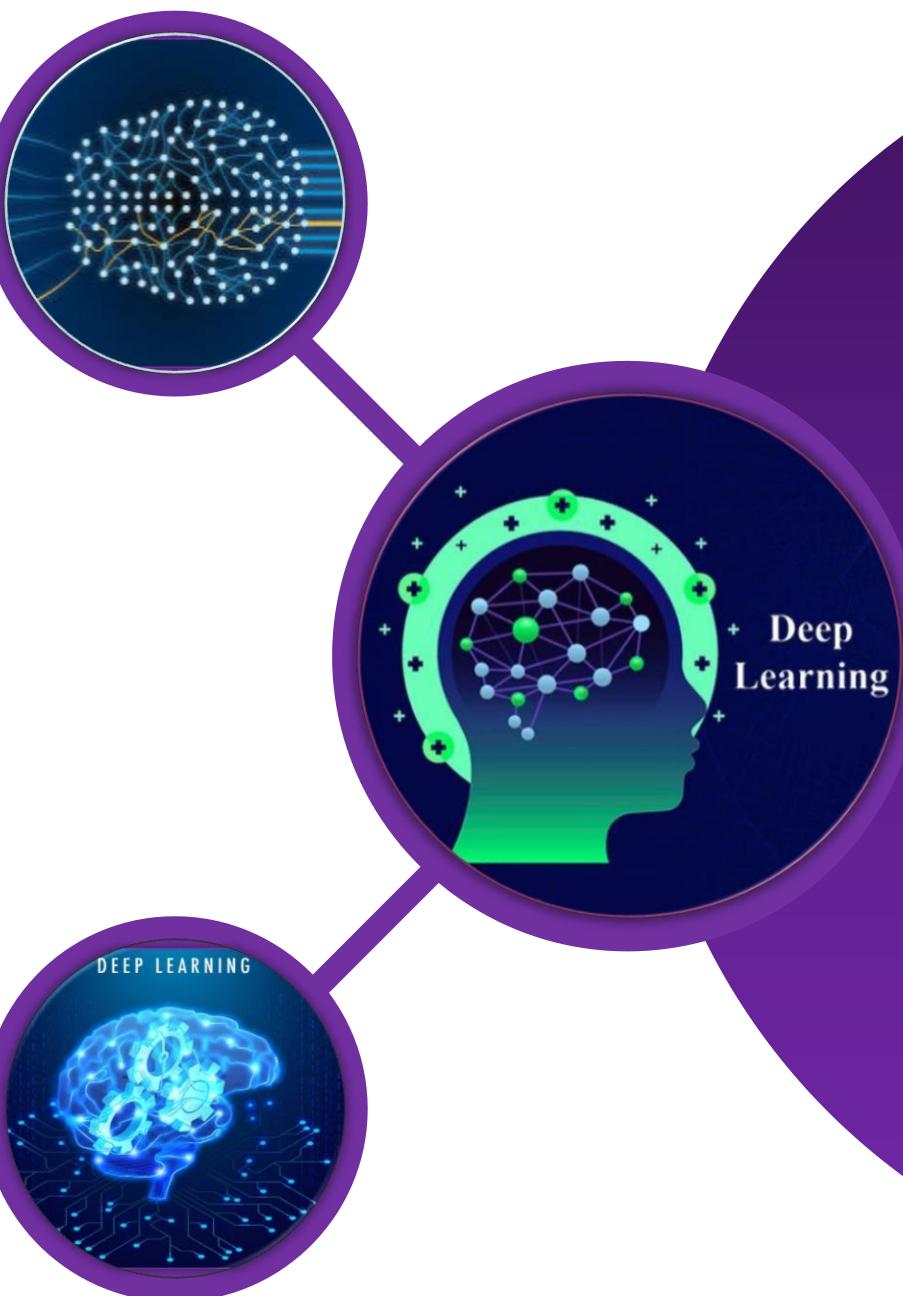
(RNN-LSTM-GRU)





PRESERVATION **CONTENT LIST**

- 01 RNN**
- 02 LSTM**
- 03 GRU**
- 04 RNN Topologies**

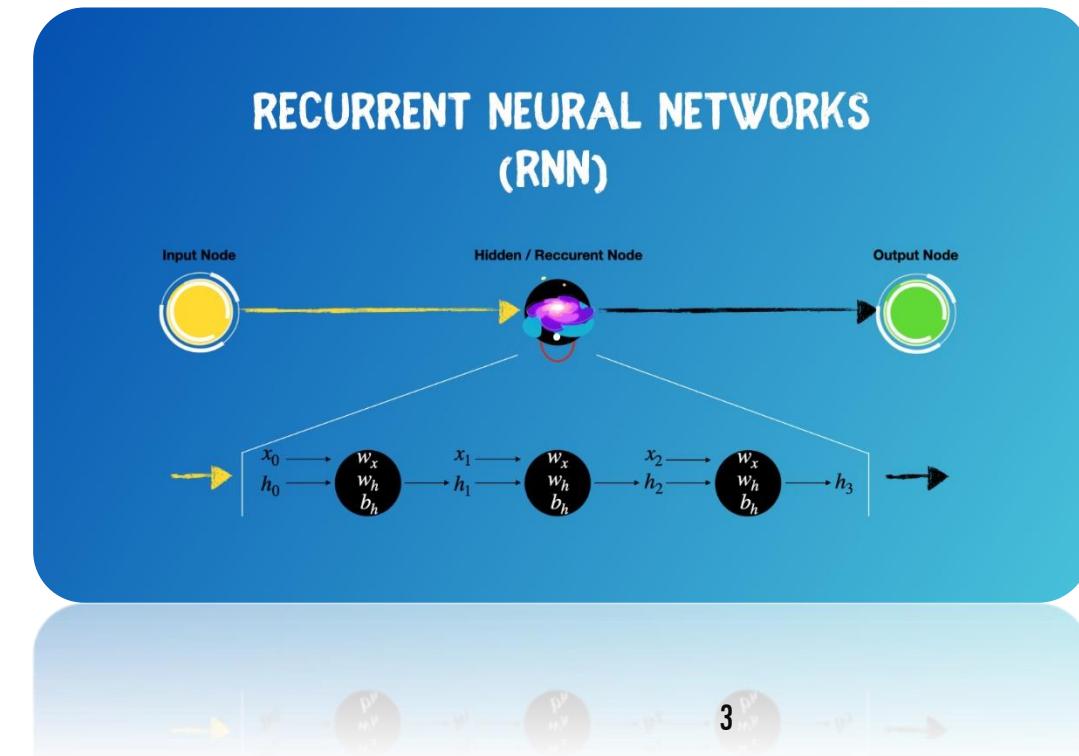
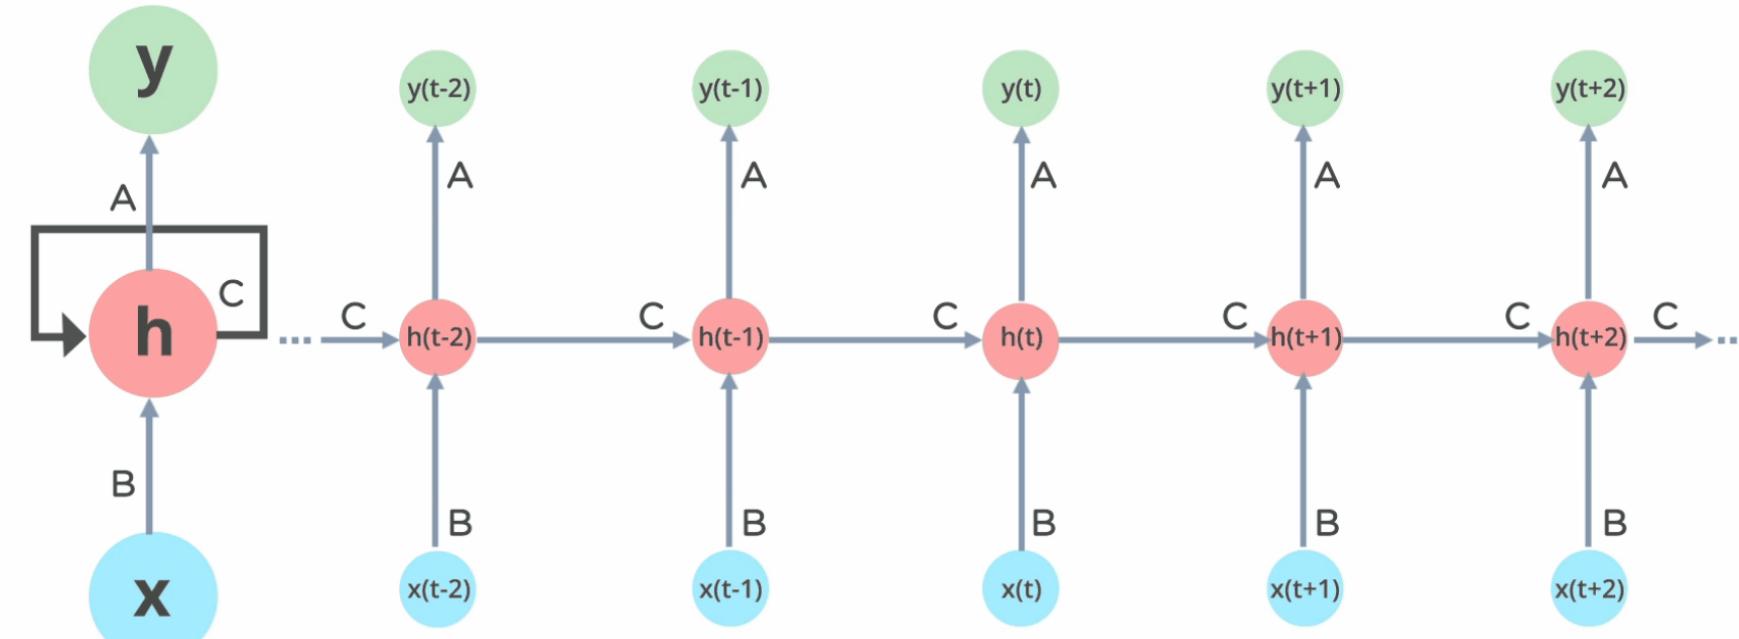


01 RNN

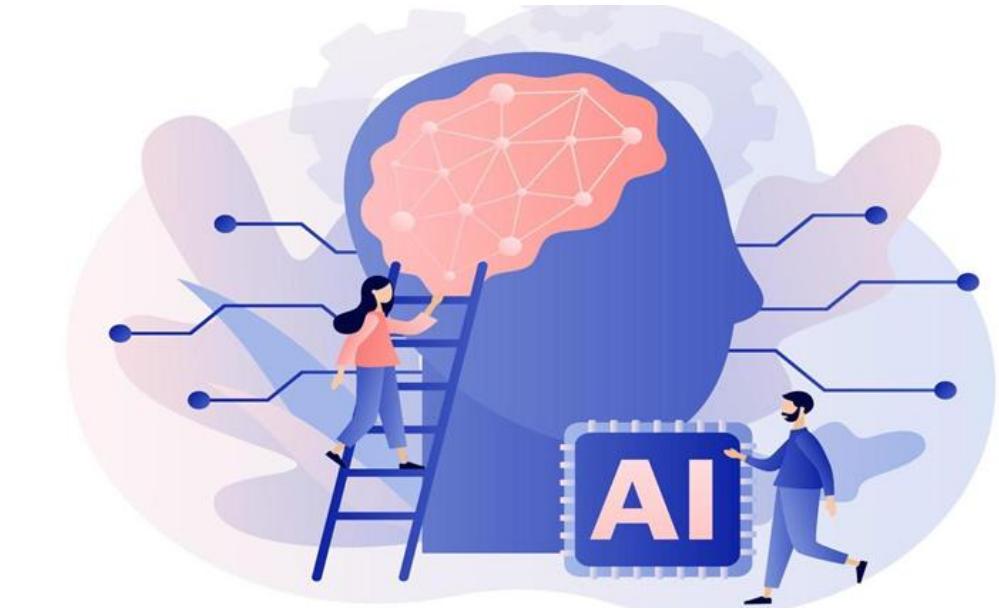


Recurrent Neural Networks (RNNs): are a family of networks that are suitable for learning representations of sequential data like text in Natural Language Processing (NLP).

- The idea behind RNNs is to make use of sequential information.
- In a traditional neural network, we assume that all inputs (and outputs) are independent of each other.
- But for many tasks that is a very bad idea. If we want to predict the next word in a sentence, we better know which words came before it.

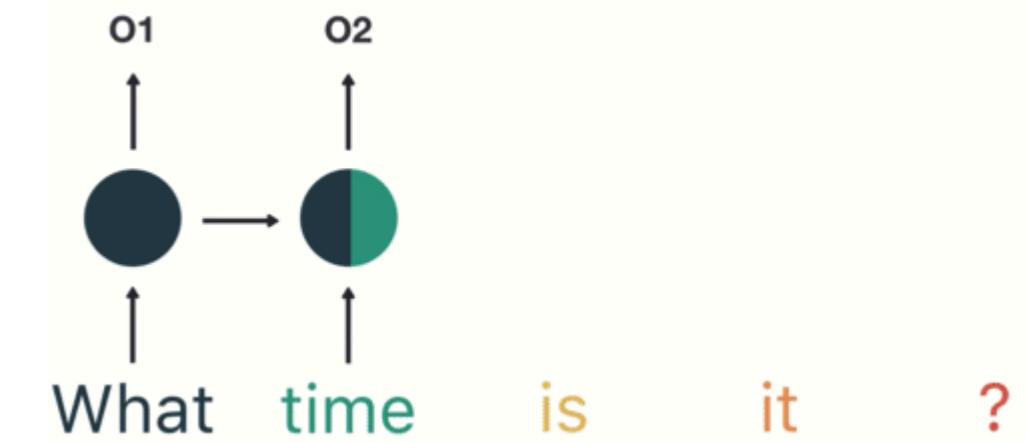
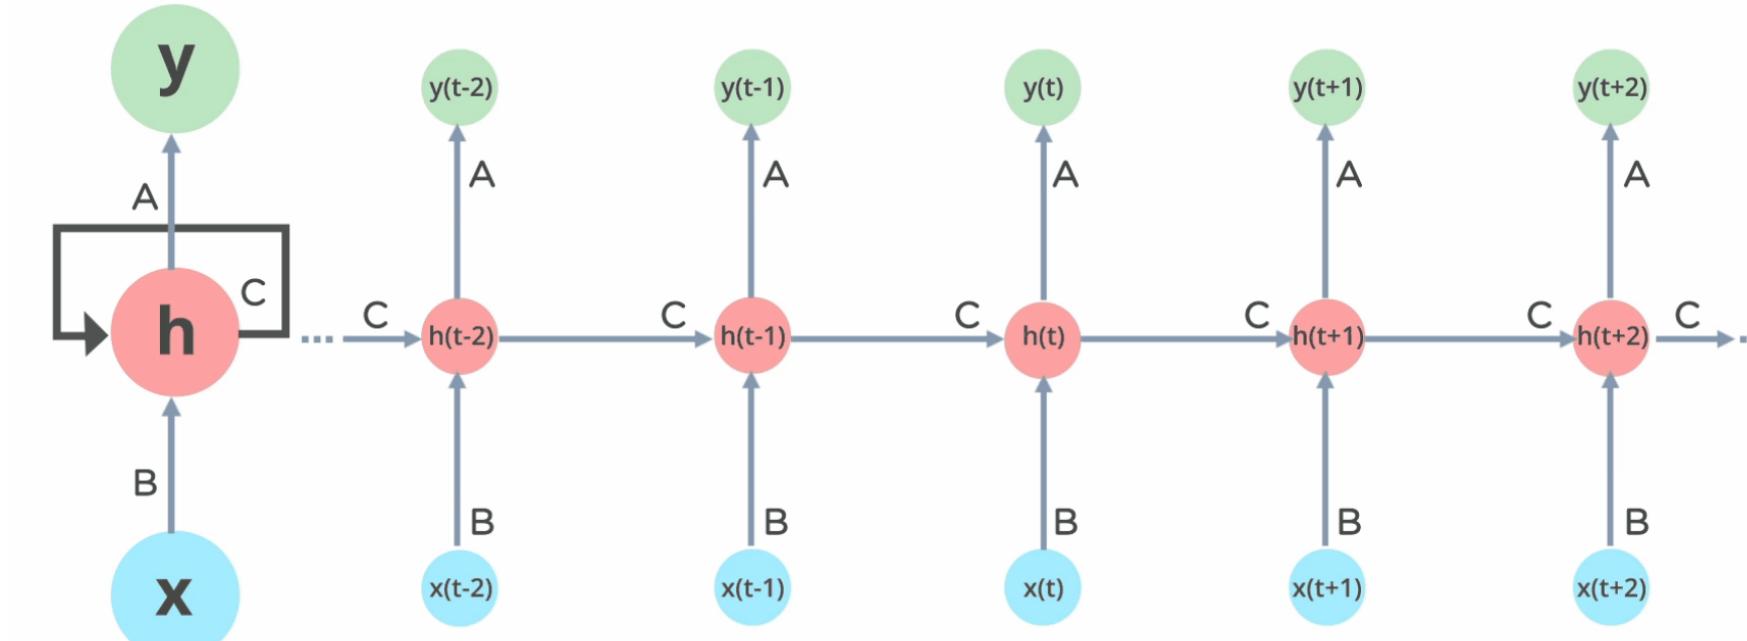


01 RNN



Recurrent Neural Networks (RNNs): are a family of networks that are suitable for learning representations of sequential data like text in Natural Language Processing (NLP).

- RNNs are called recurrent because they perform the same task for every element of a sequence, with the output being depended on the previous computations.
- Another way to think about RNNs is that they have a “memory” which captures information about what has been calculated so far.
- In theory RNNs can make use of information in arbitrarily long sequences, but in practice they are limited to looking back only a few steps.

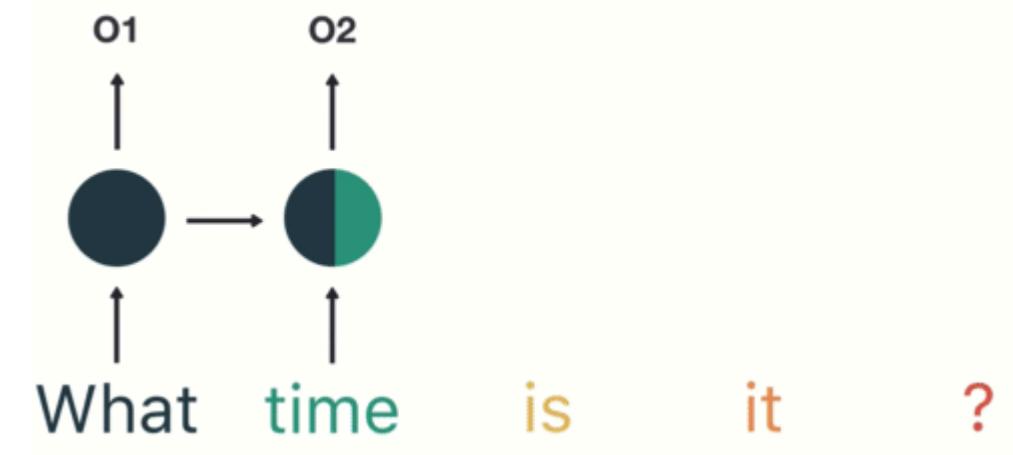
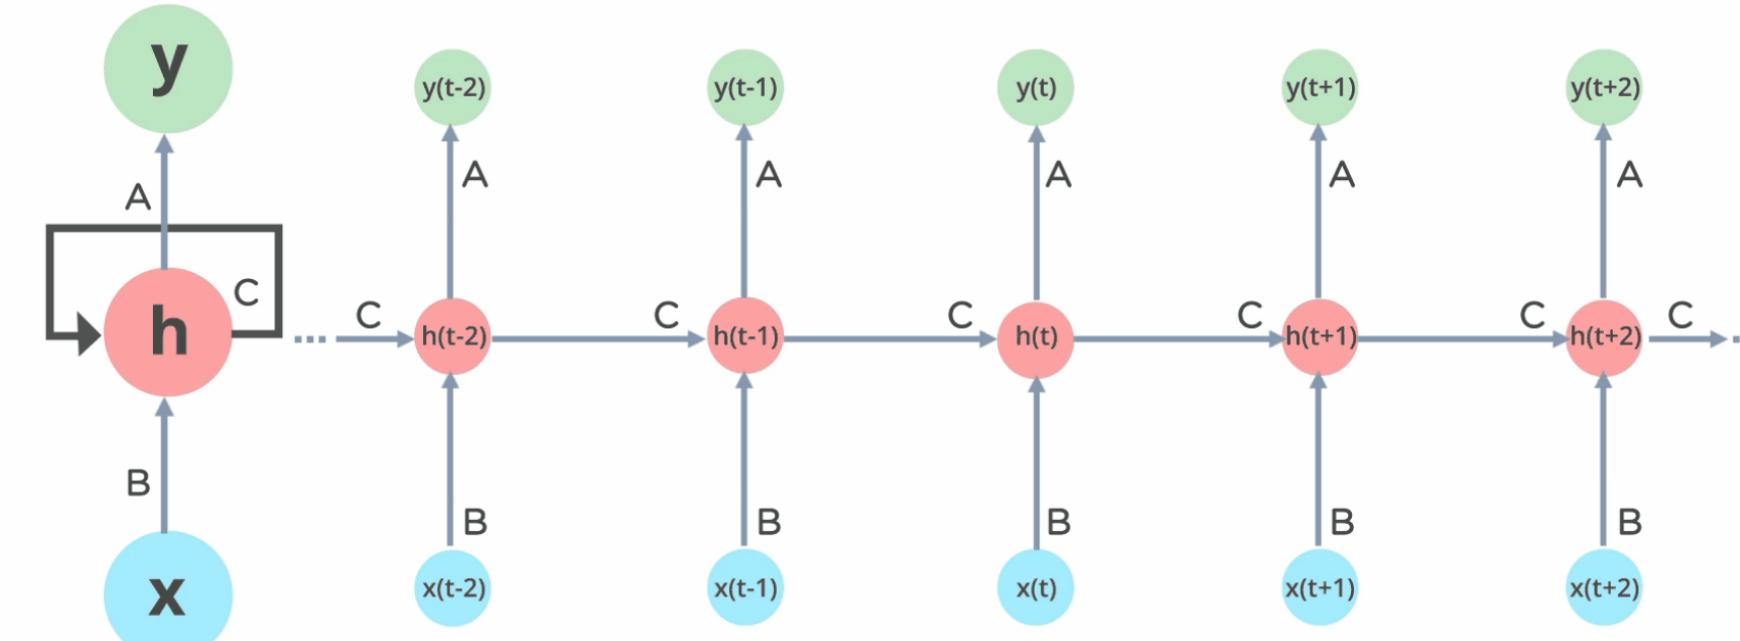


01 RNN

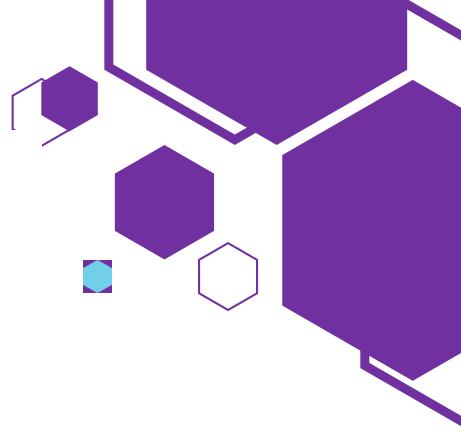
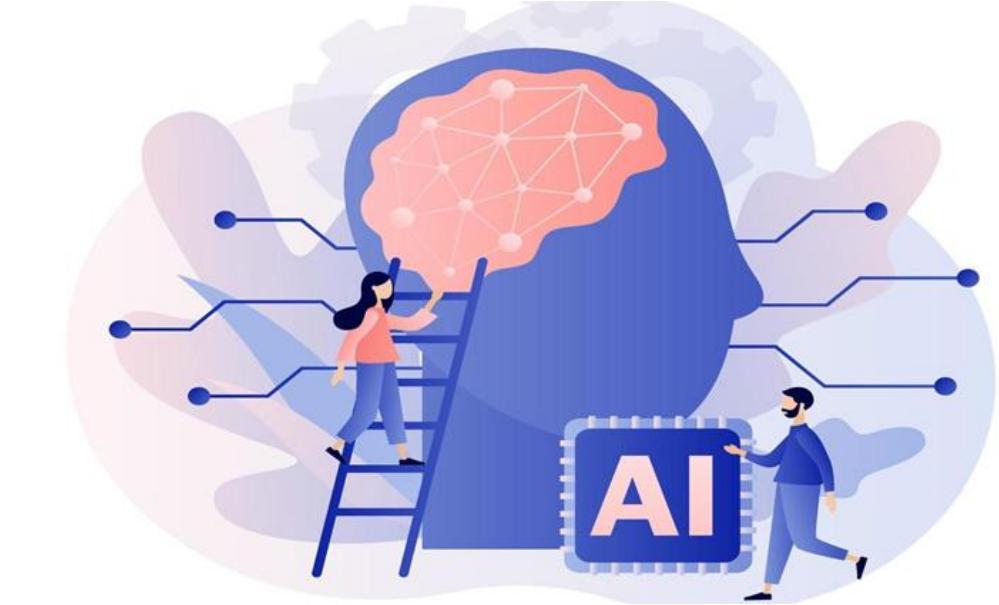


Recurrent Neural Networks (RNNs): are a family of networks that are suitable for learning representations of sequential data like text in Natural Language Processing (NLP).

- The above diagram shows a RNN being unrolled (or unfolded) into a full network.
- For example, if the sequence we care about is a sentence of 5 words, the network would be unrolled into a 5-layer neural network, one layer for each word.

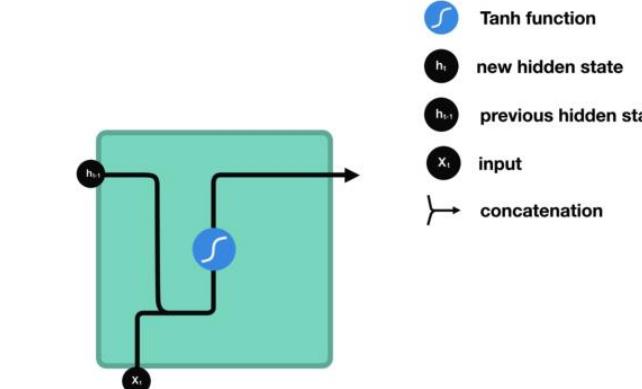


01 RNN

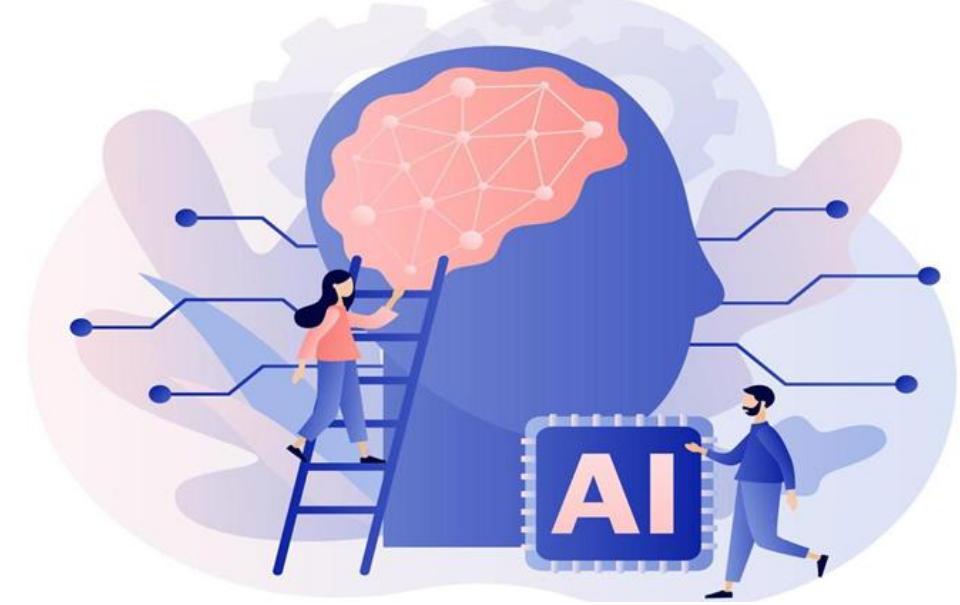


Recurrent Neural Networks (RNNs): are a family of networks that are suitable for learning representations of sequential data like text in Natural Language Processing (NLP).

- Time series data, such as stock prices, also exhibit a dependence on past data, called the secular trend.
- RNN cells incorporate this dependence by having a hidden state, or memory, that holds the essence of the past.
- The value of the hidden state at any point in time is a function of the value of the hidden state at the previous time step and the value of the input at the current time step, that is: $h_t = \tanh(W * h_{t-1} + U * X_t)$
- h_t and h_{t-1} are the values of the hidden states at the time steps t and t-1 respectively, and x_t is the value of the input at time t.
- The above equation is recursive, that is, h_{t-1} can be represented in terms of h_{t-2} and x_{t-1} , and so on, until the beginning of the sequence.



01 Vanishing gradient



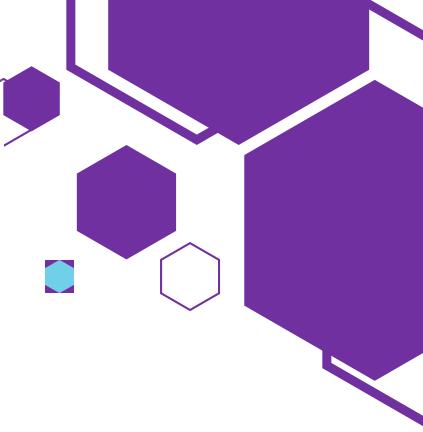
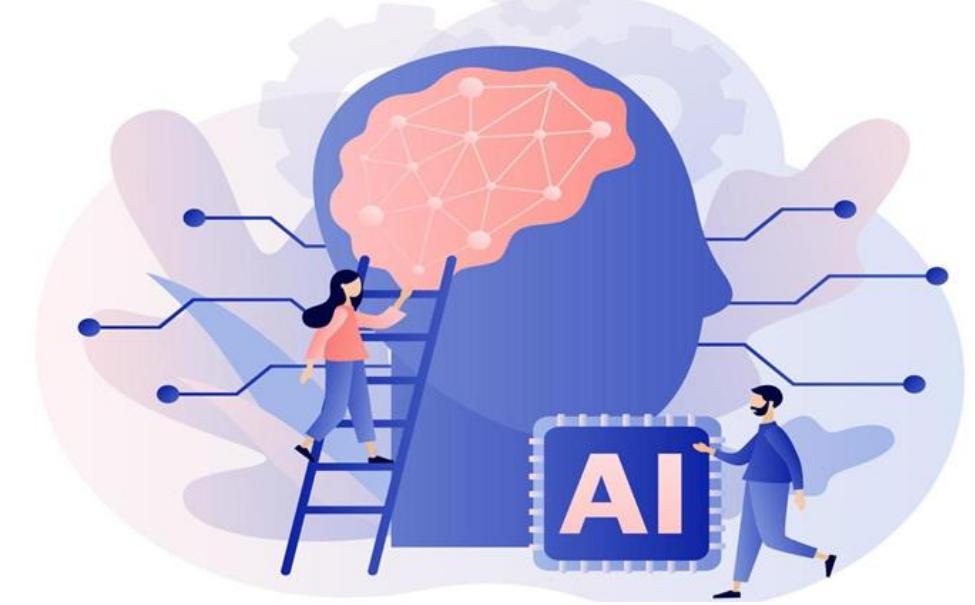
When doing back propagation, each node in a layer calculates its gradient with respect to the effects of the gradients in the layer before it.

- Just like traditional neural networks, training the RNN also involves backpropagation.
- The difference in this case is that since the parameters are shared by all time steps, the gradient at each output depends not only on the current time step, but also on the previous ones.
- This process is called backpropagation through time (BPTT).
- Regular RNNs might have a difficulty in learning long range dependencies.
- This kind of dependencies between sequence data is called long-term dependencies because the distance between the relevant information and the point where it is needed to make a prediction is very wide.

- **RNNs can't learn long-range dependencies!**

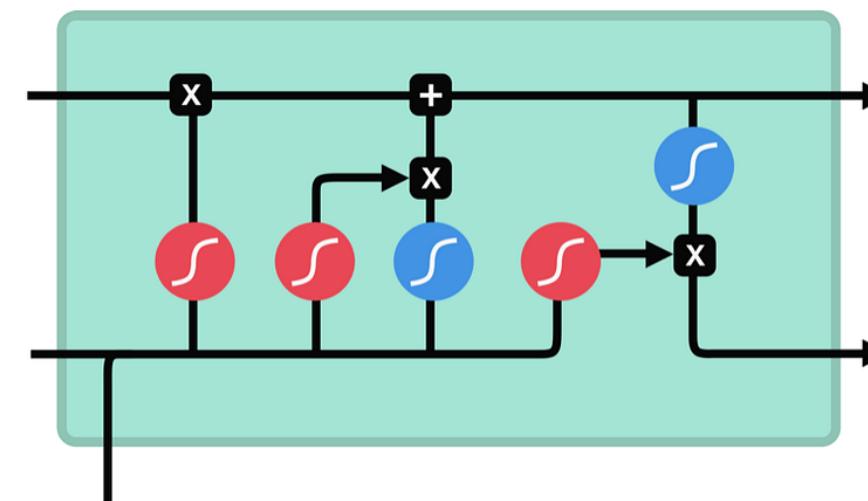


02 LSTM



Similar in nature to RNNs but have additional features to help fight the short-term memory issue of RNNs.

While there are a few approaches to minimize the problem of vanishing gradients, such as proper initialization of the W matrix, using a ReLU instead of tanh layers, and pre-training the layers using unsupervised methods, the most popular solution is to use the LSTM or GRU architectures.



sigmoid

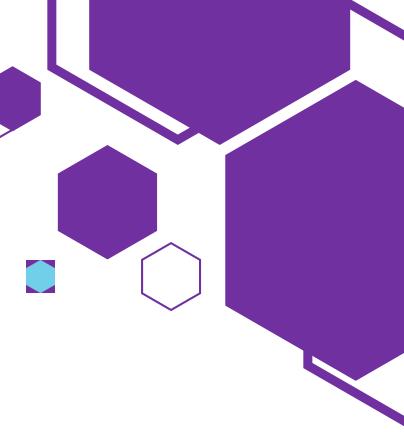
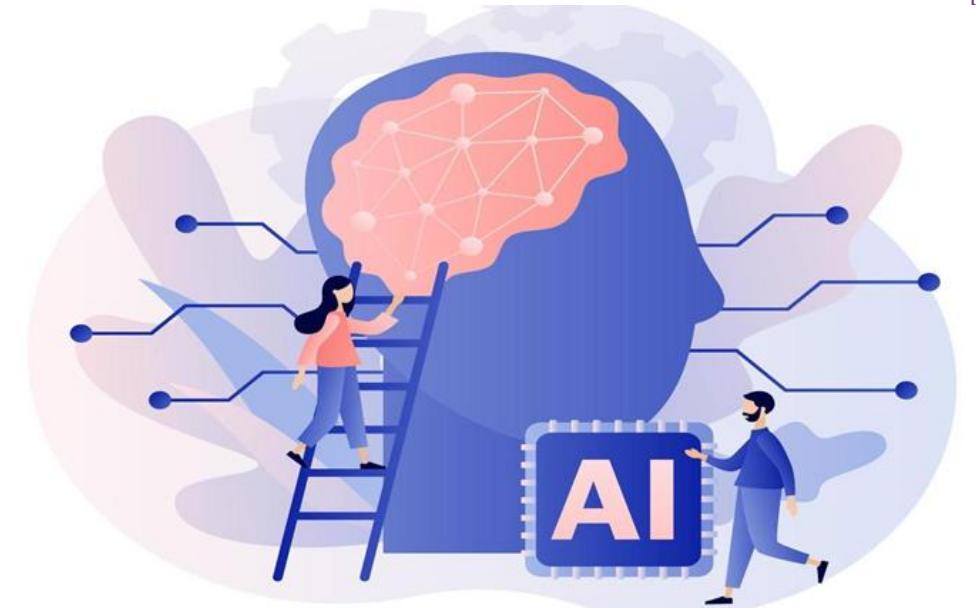
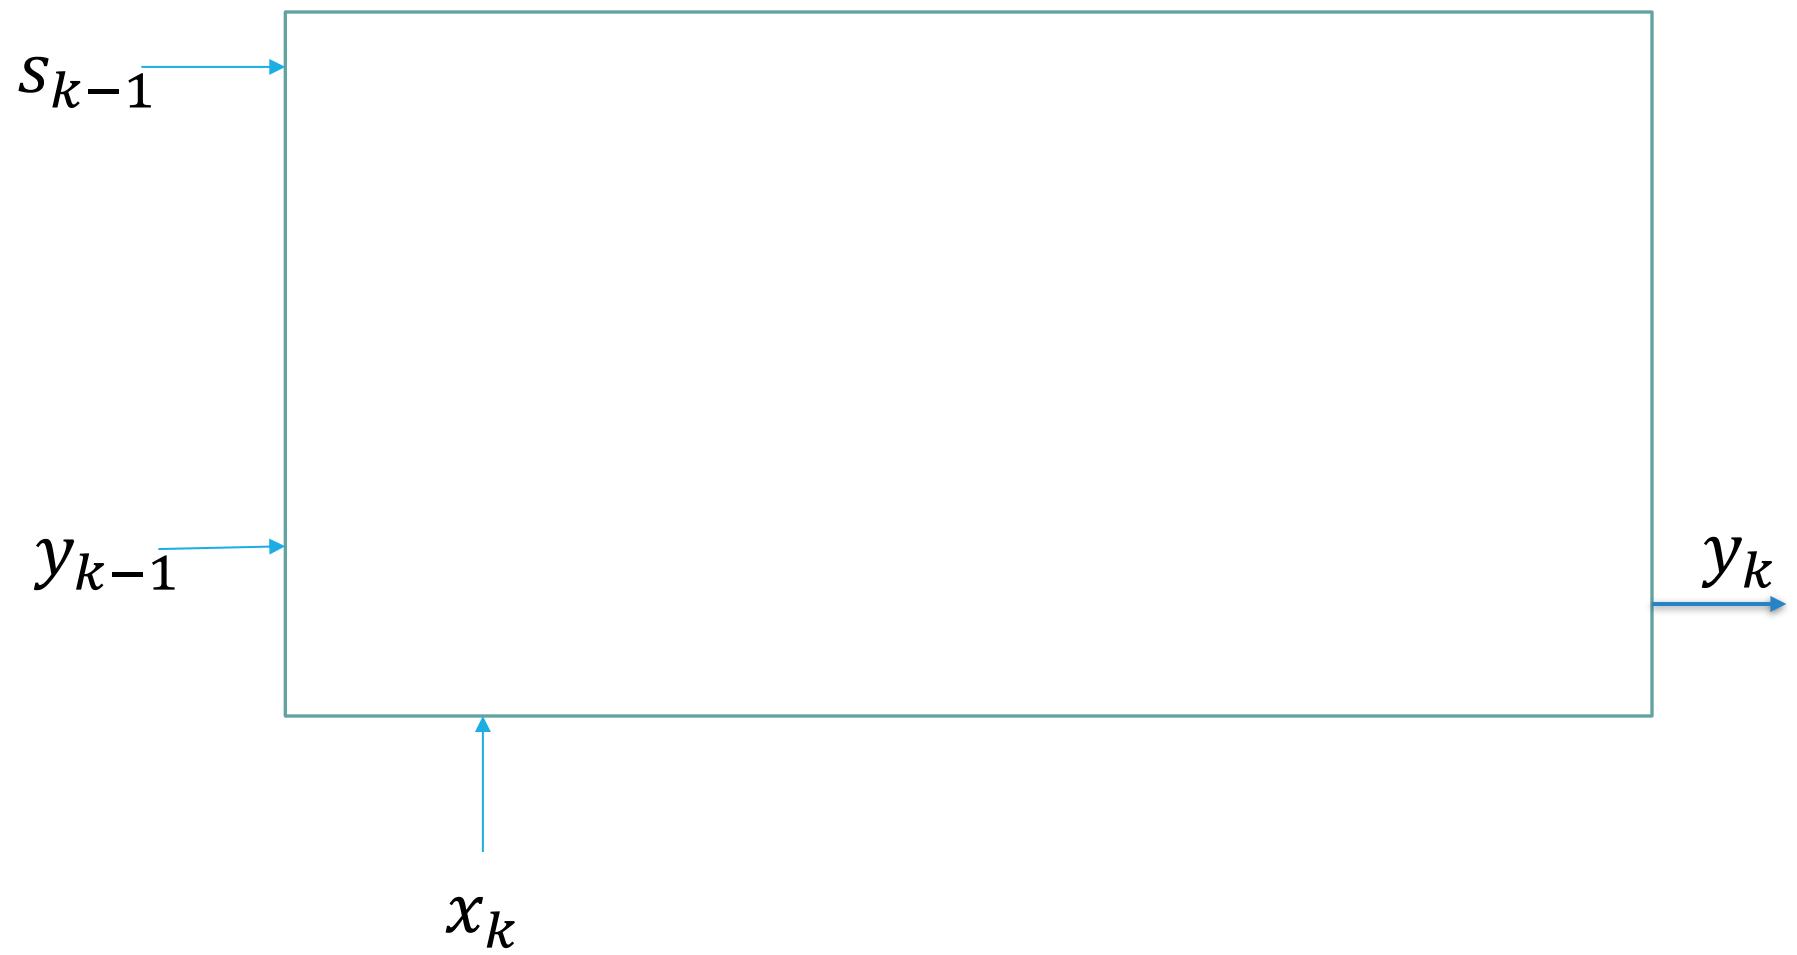
tanh

pointwise
multiplication

pointwise
addition

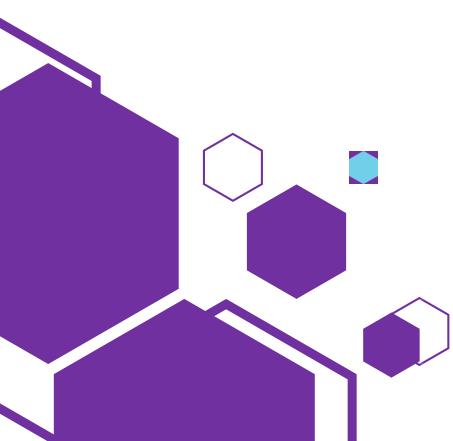
vector
concatenation

02 LSTM

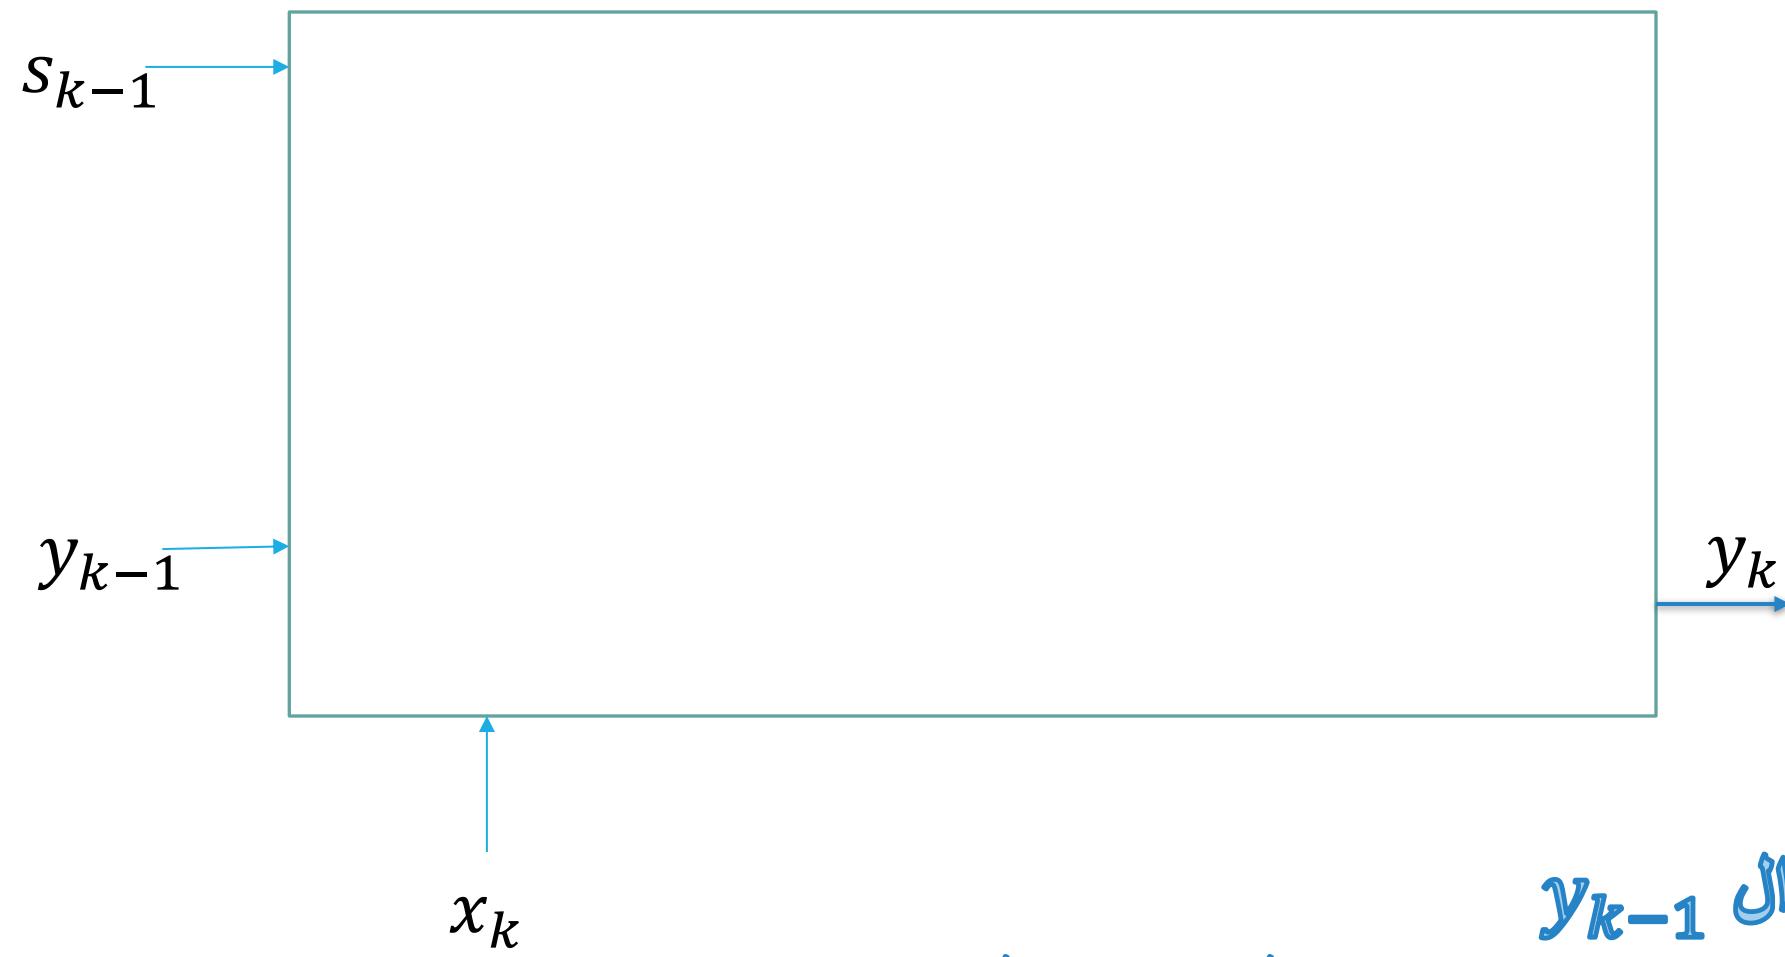
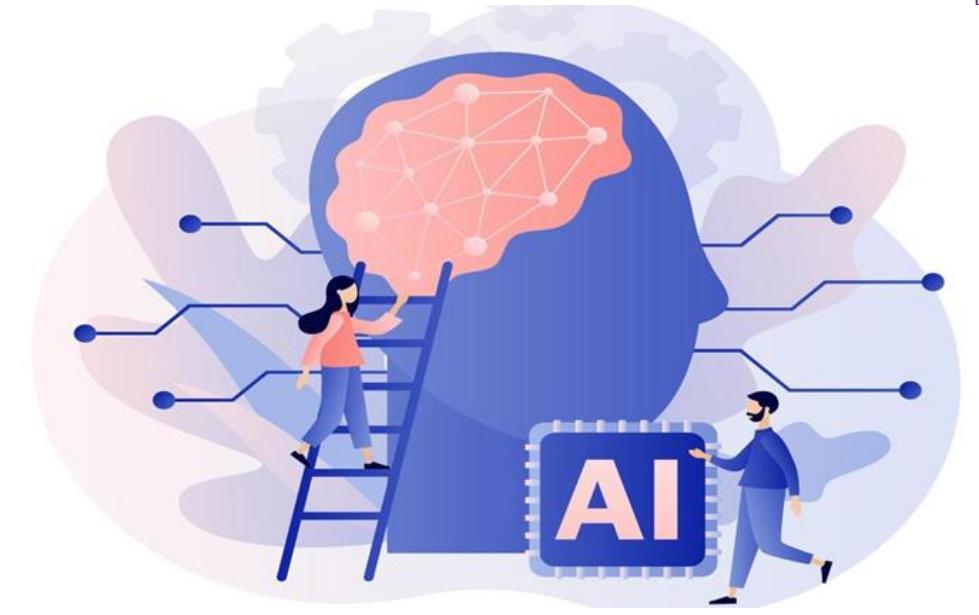


انا دلوقتى عايز اخذ قرار معين y_k
القرار ده هيعتمد على 3 حاجات:

- 1- الانبوت اللي جايلى x_k اللي أنا شايفه دلوقتى
- 2- اخر قرار أنا اخذته y_{k-1}
- 3- وكل القرارات اللي اخذتها قبل كده ال history بتعنى يعني s_{k-1}

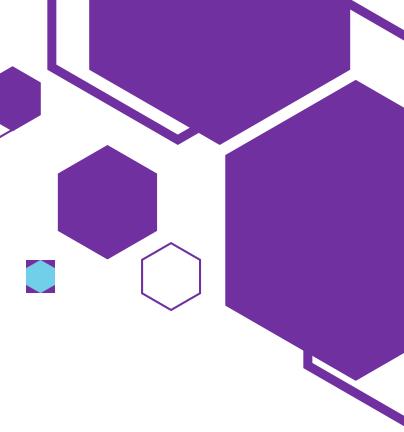
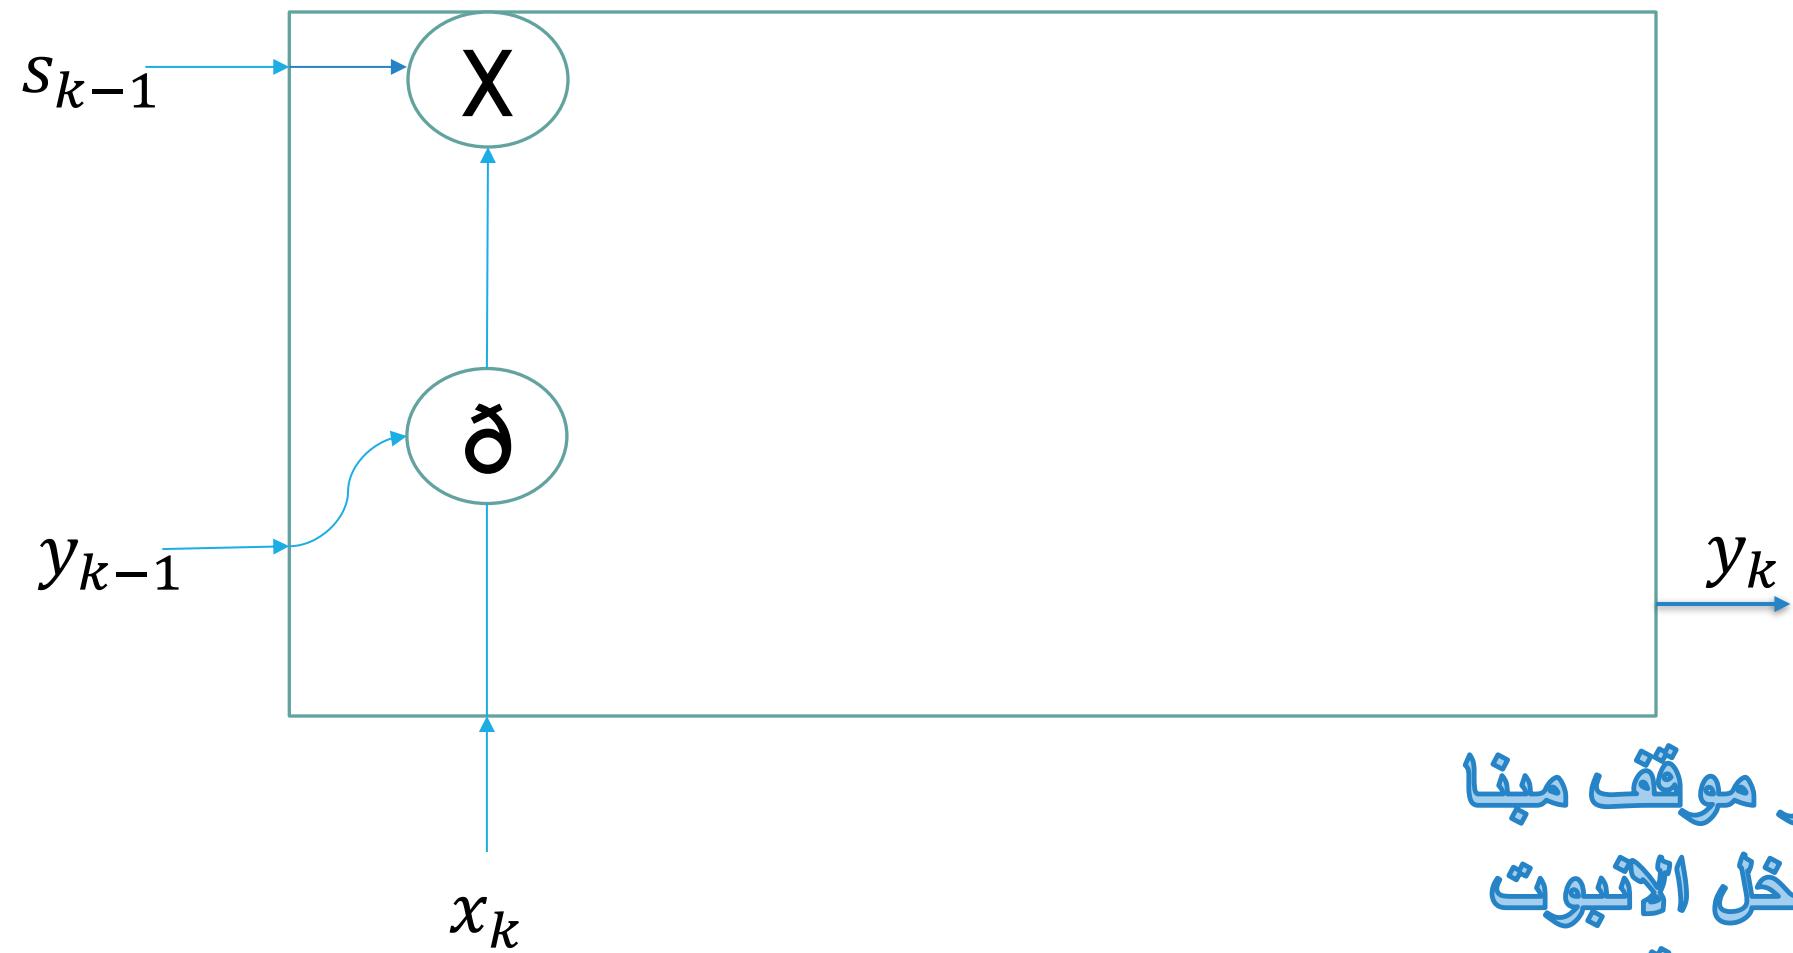


02 LSTM

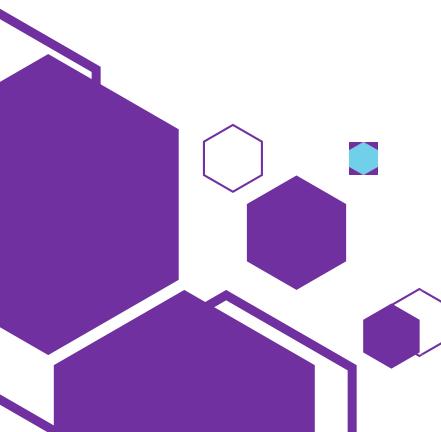


- 1- ولیکن قولتک فلان بیقول علیک کلام وحش بیقی ده ال x_k
- 2- فلان ده انا آخر موقف کان معاه کان ایه موقف رجوله ولا بیع ده ال y_{k-1}
- 3- طب تعالی بقی نرجع بالذکر ونقلب فی الماضي كل الی فات میینی ومیینوا کان عامل ایه ده ال s_{k-1} بیقی هتربط الموقف الی حصل دلوقتی باخر موقف حصل وبكل المواقف القديمه الی مبنکوا من يوم معرفته وعلى أساسه هاخد قرار الی هو مثلًا عفا الله عما سلف ولا اقطع علاقتي بيه ولا احذو من التعامل پس معاه ولا ایه

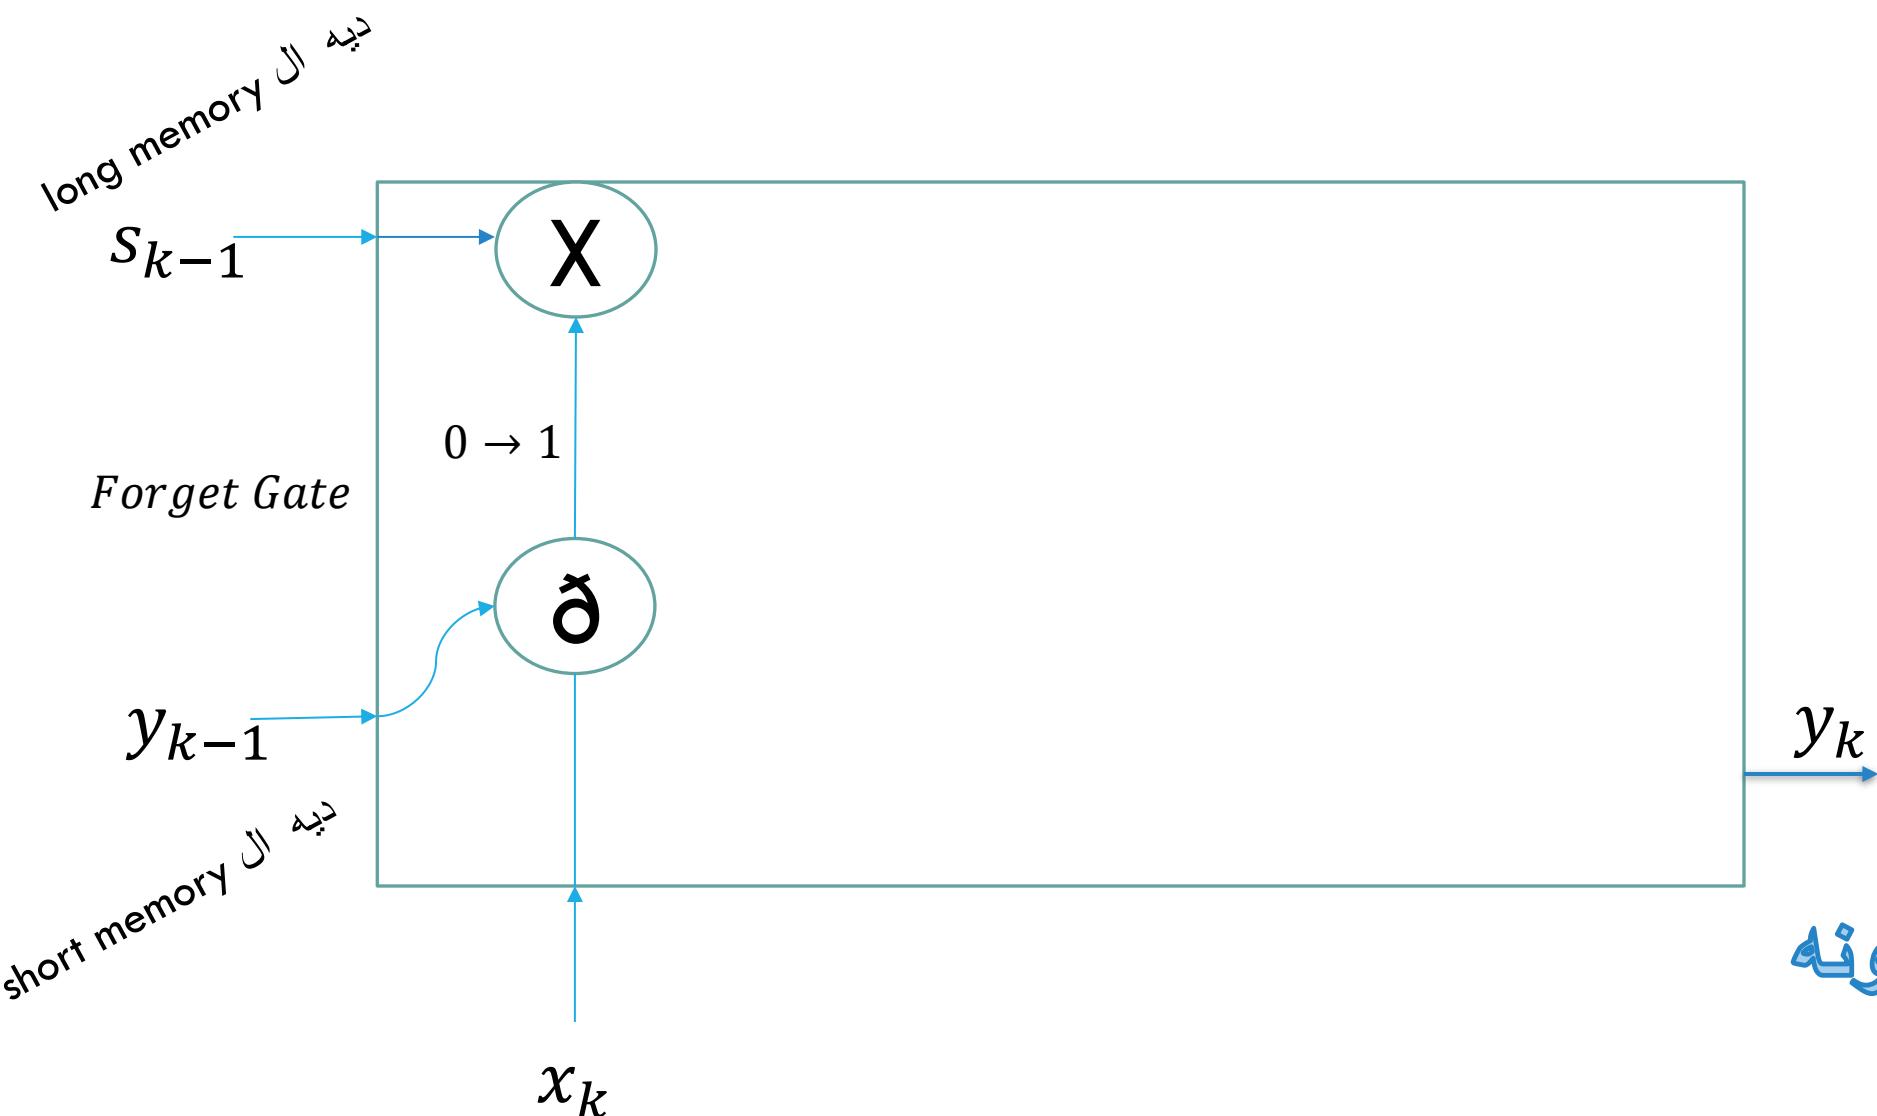
02 LSTM



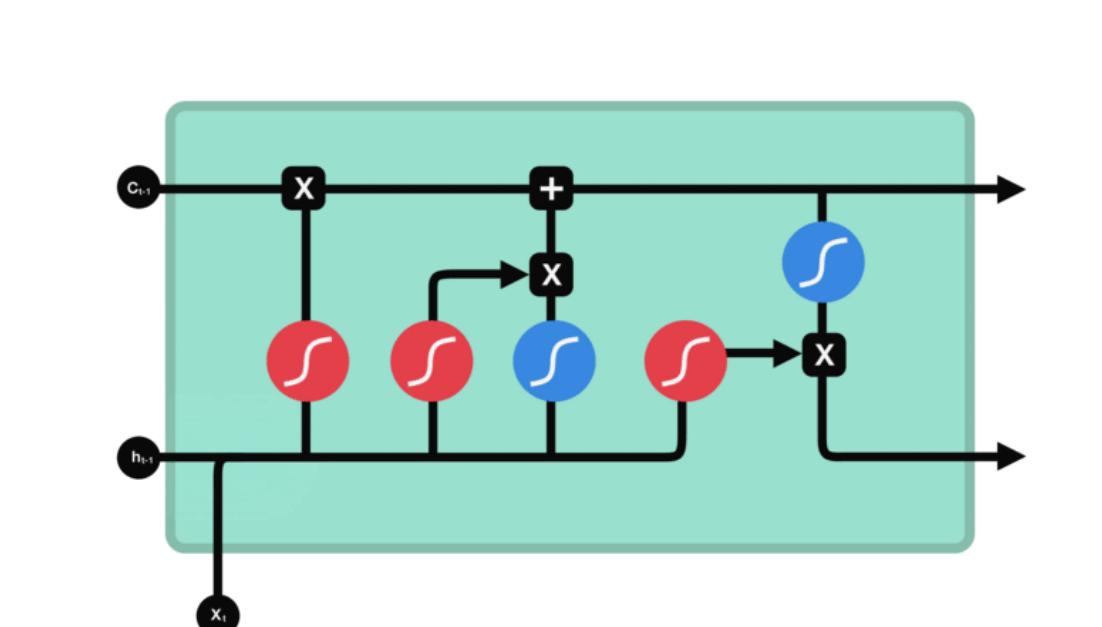
اڪٽر حاجه هيڪى ليها تأثير على قرارى هيڪى الكلام اللى جايلى واخر موقف مينا انما التاريخ الأسود اللى مينا مش هيڪى ليه تأثير قوى عشان كده هدخل الانبوت والآوت السابق على sigmoid بحبيث اشوف التاريخ الأسود هيڪر بنسبة قد ايه مش يمڪن الموقف اللى حصل واخر موقف مينا يعنى انى اشوف التاريخ فالاوتبوت بتاع ال sigmoid وقتها هيڪى بصفر فمش هبص عاليهستوري واقطع علاقتى بيء



02 LSTM



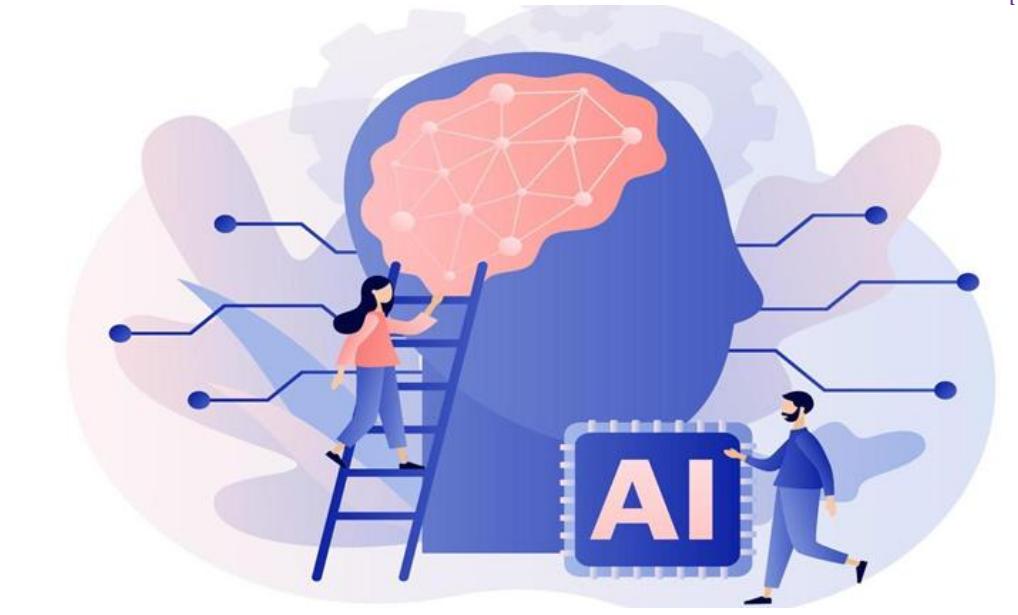
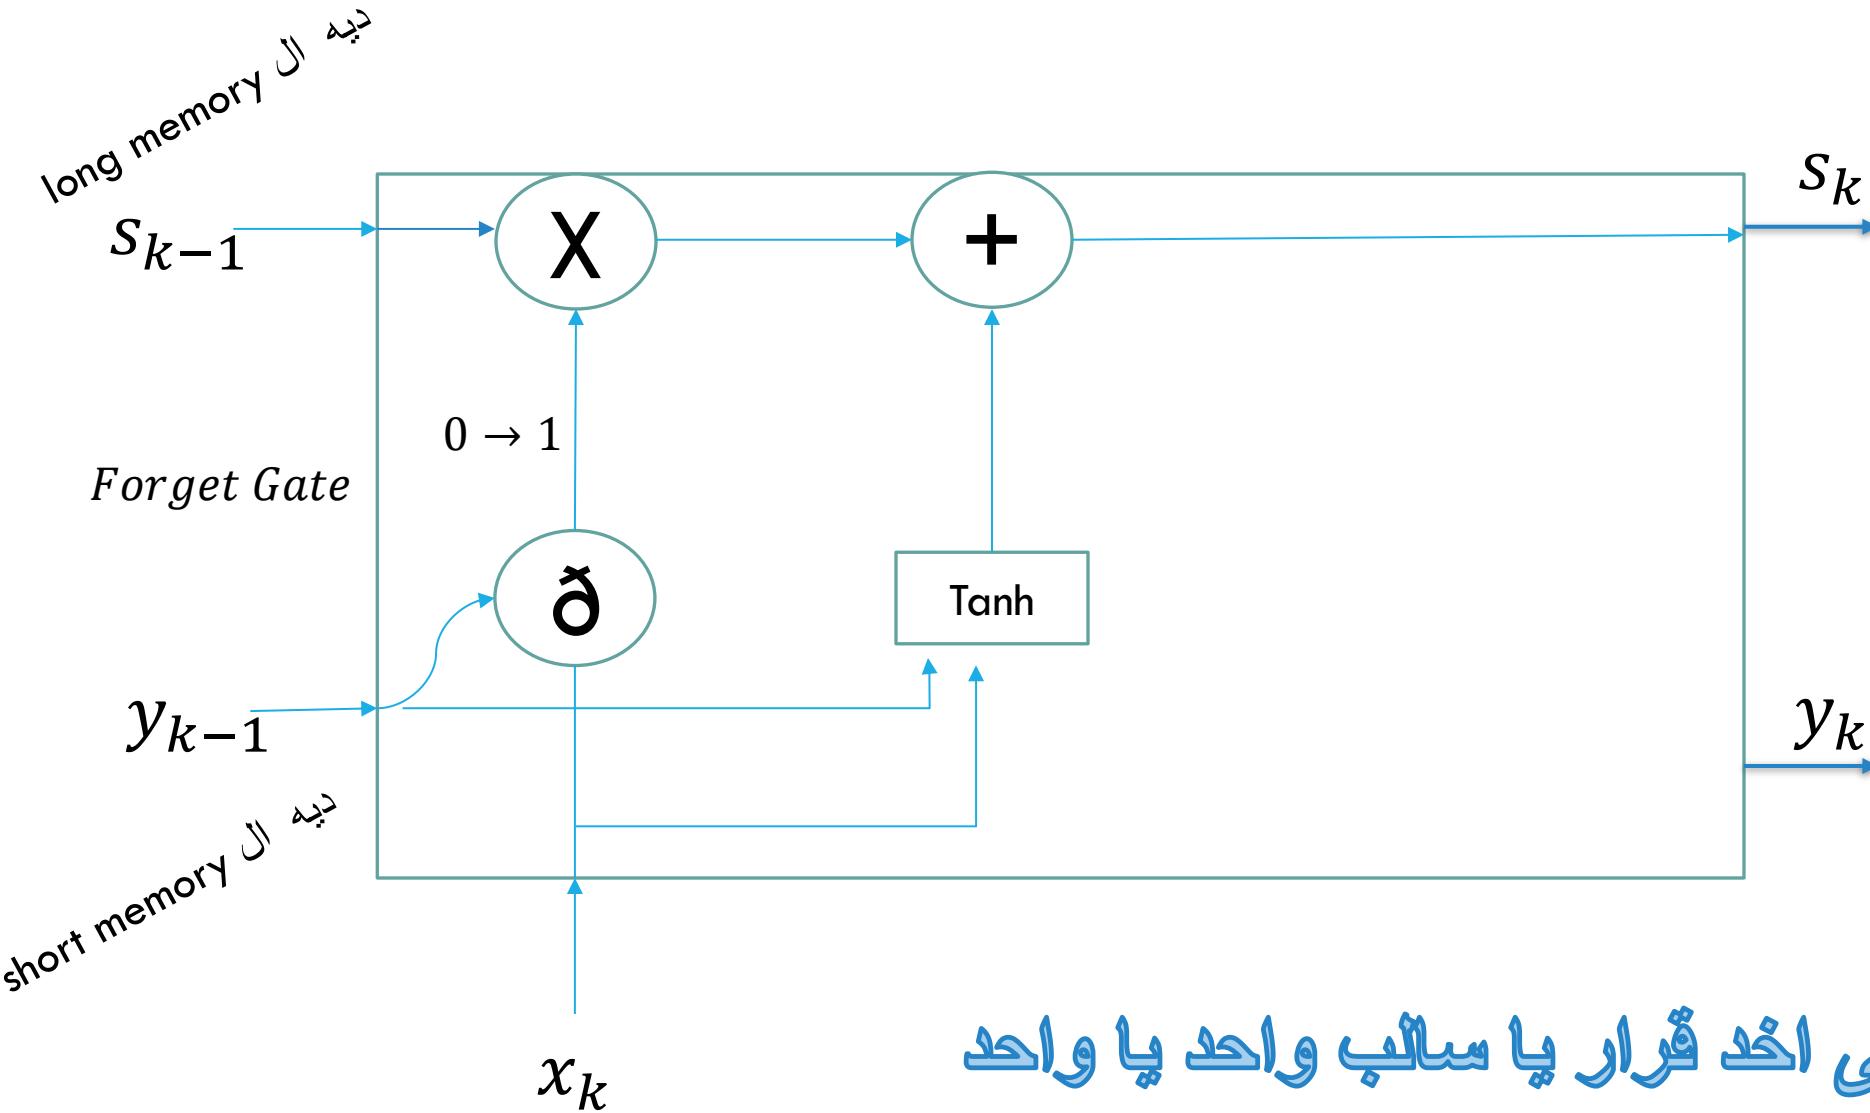
بما ان دیه بتعتمد على اني هاخذ قد ايه من الذكريات وال حاجات المدفونه
فهيسمىها forget gate



- \bullet c_{t-1} previous cell state
- \bullet f_t forget gate output

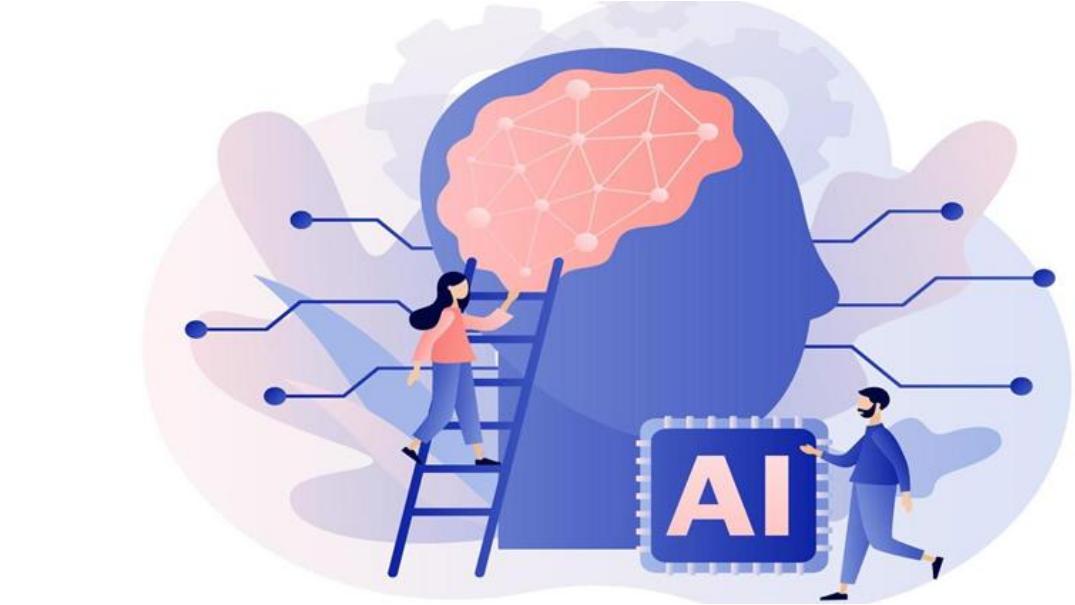
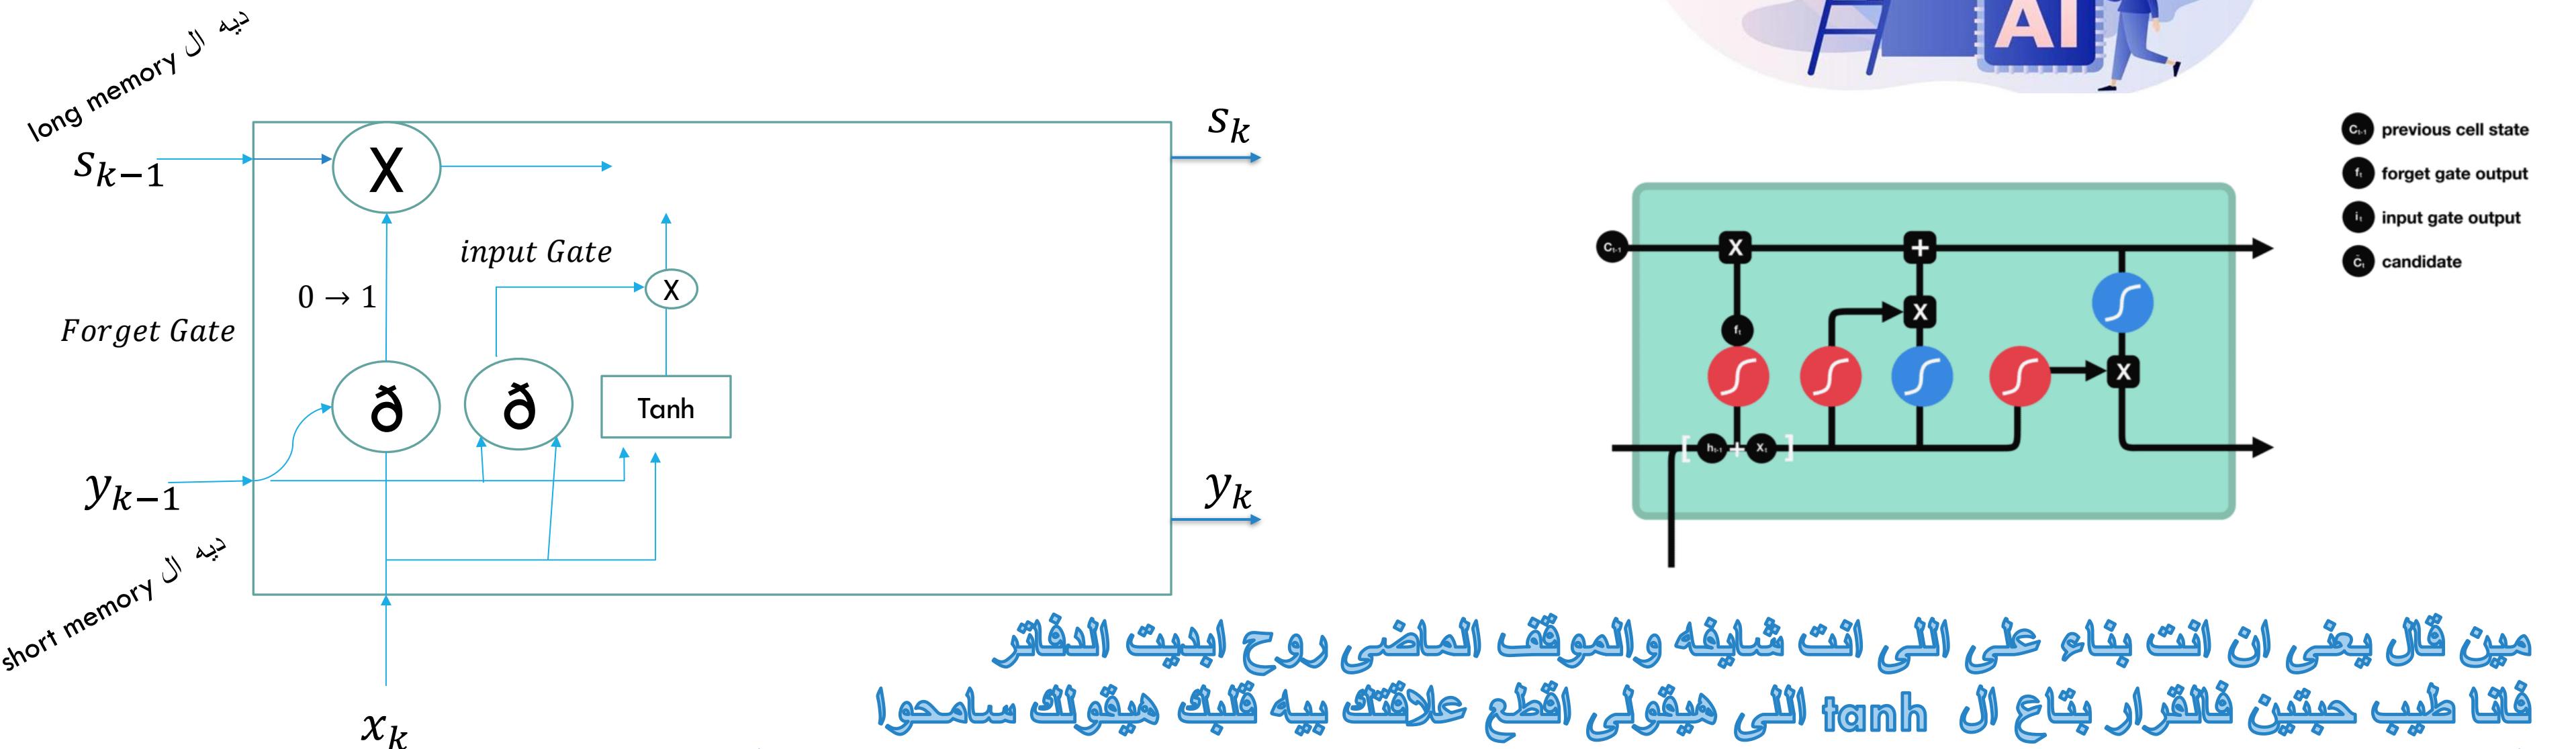
- **Forget gate:** Information from the previous hidden state and information from the current input is passed through the sigmoid activation function.
- Values come out between 0 and 1
- Decides what information should be thrown away or kept

02 LSTM

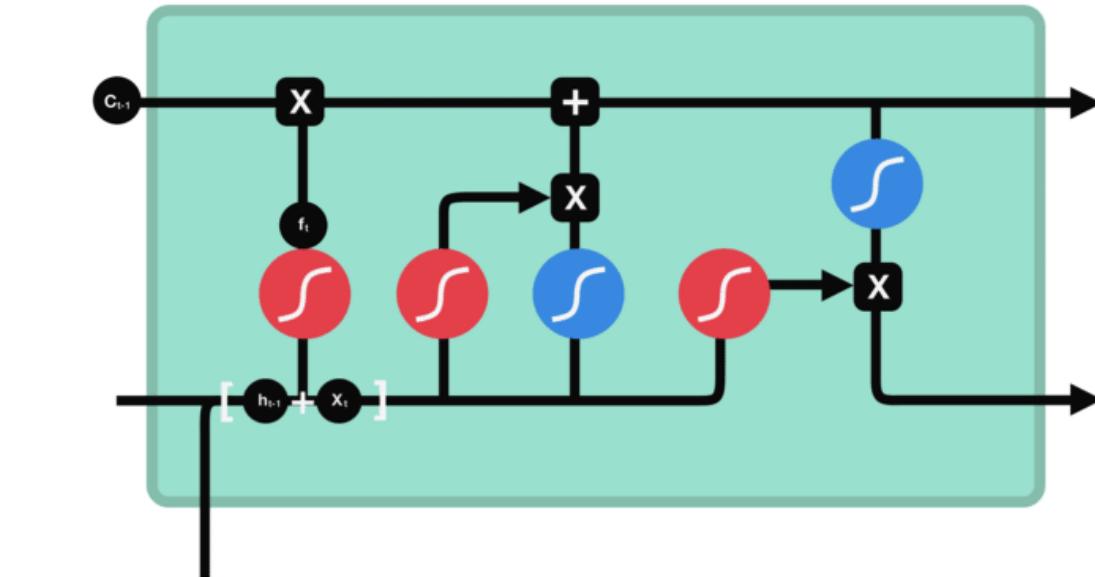


دلوقتی عایز أكون ال state الجديد بمعنی بابیت الدفاتر القديمه
فهآخذ الانبوت والاوتيوت القديم وادخلهم على tanh ليه اللي هتخلينى اخذ قرار با سالب واحد با واحد
فهآخذ القرار ده اجمعوا على الاوتيوت اللي جايلى من ال history

02 LSTM



- c_{t-1} previous cell state
- f_t forget gate output
- i_t input gate output
- \tilde{c}_t candidate

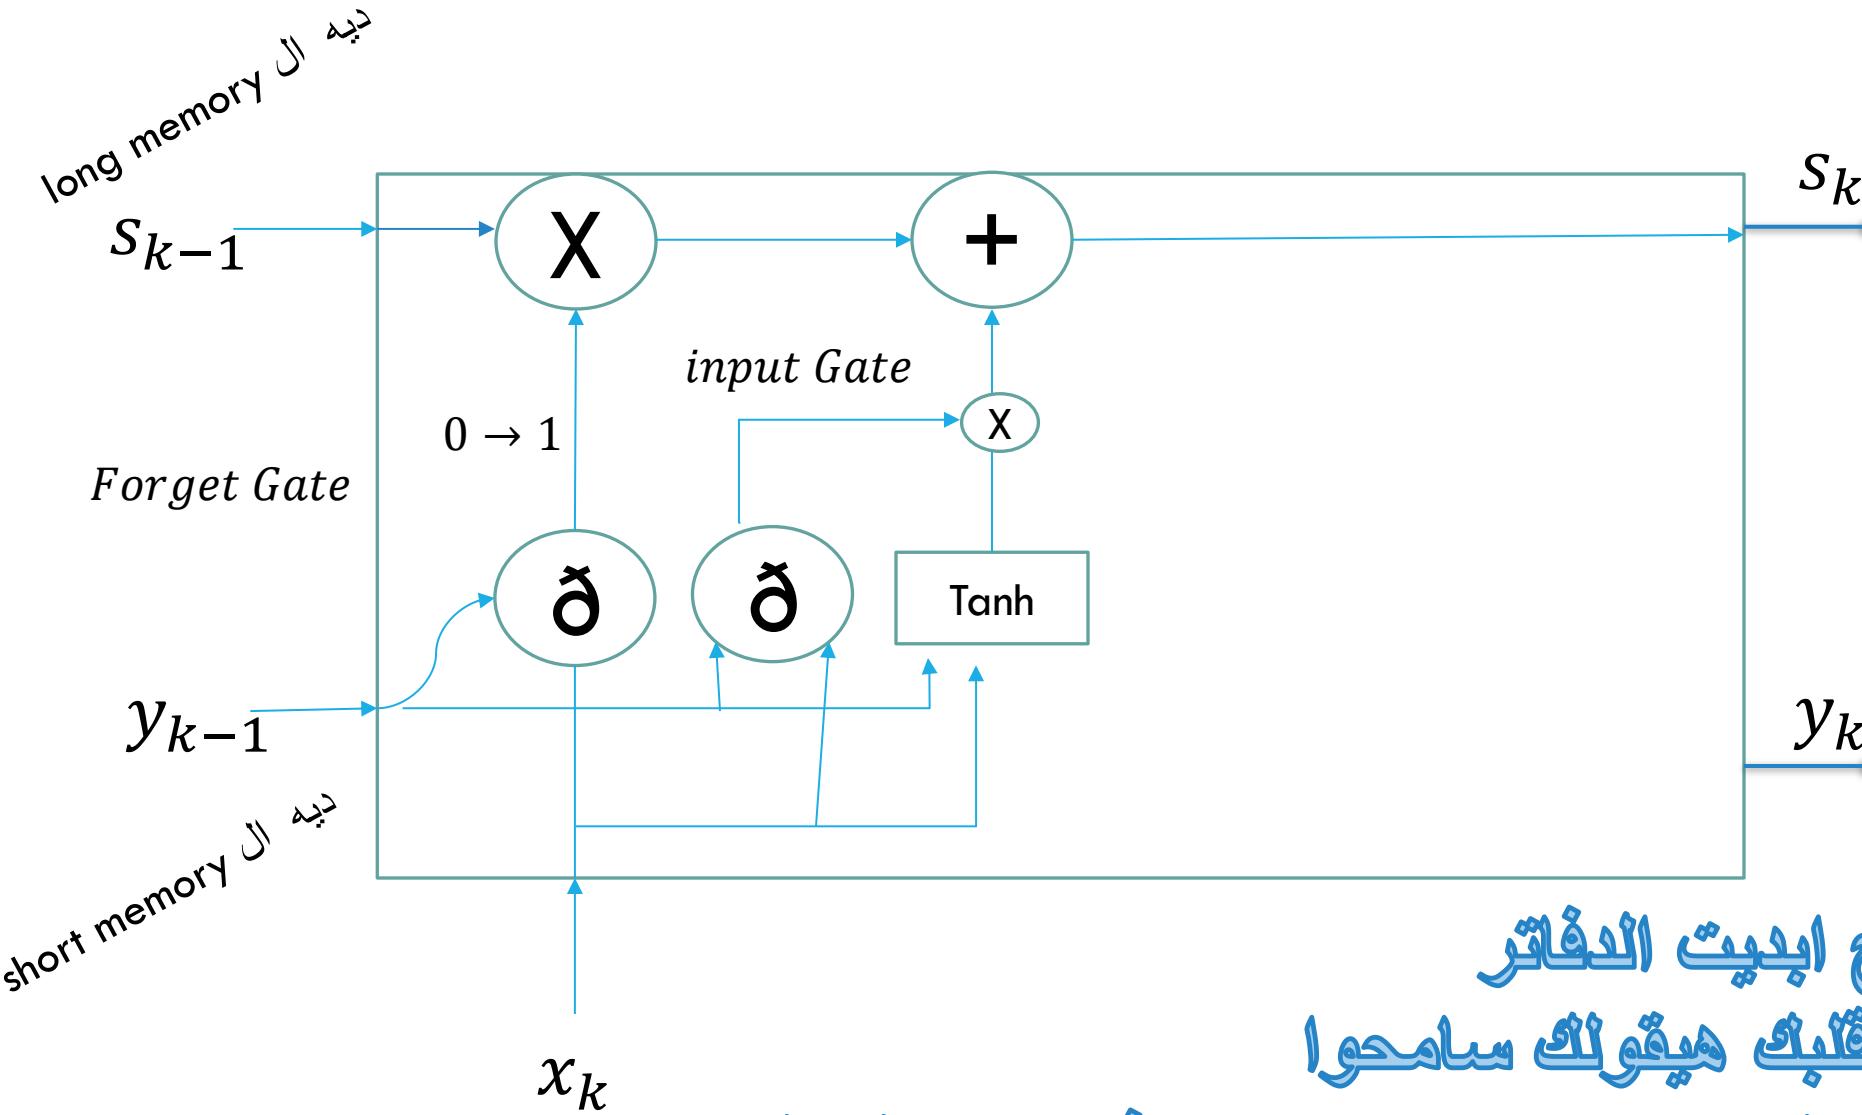


مین قال یعنی ان انت بناء على اللى انت شايفه والموقف الماضى روح ابديت الدفاتر
فانا طيب جيتين فالقرار بتاع ال $tanh$ اللى هيقولى اقطع علاقتك بيه قلبك هيقولك سامحوا
فهمخذ الانبوت الحالى والانبوت السابق وادخله على $sigmoid$ جديده اخلى الاوت بتاعها يضرب في اوت التانش
فيه هسيبيا $input gate$ اللى هو انبوت جديد للتاريخ بتاعي

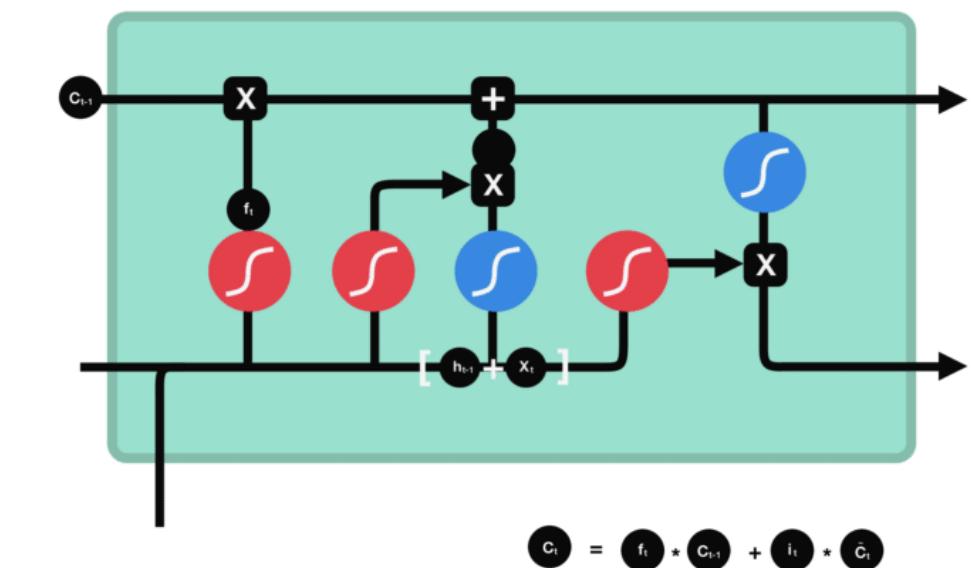
- Input gate

Decides what new information we're going to store in the cell state

02 LSTM



- c_{t-1} previous cell state
- f_t forget gate output
- i_t input gate output
- \hat{c}_t candidate
- c_t new cell state

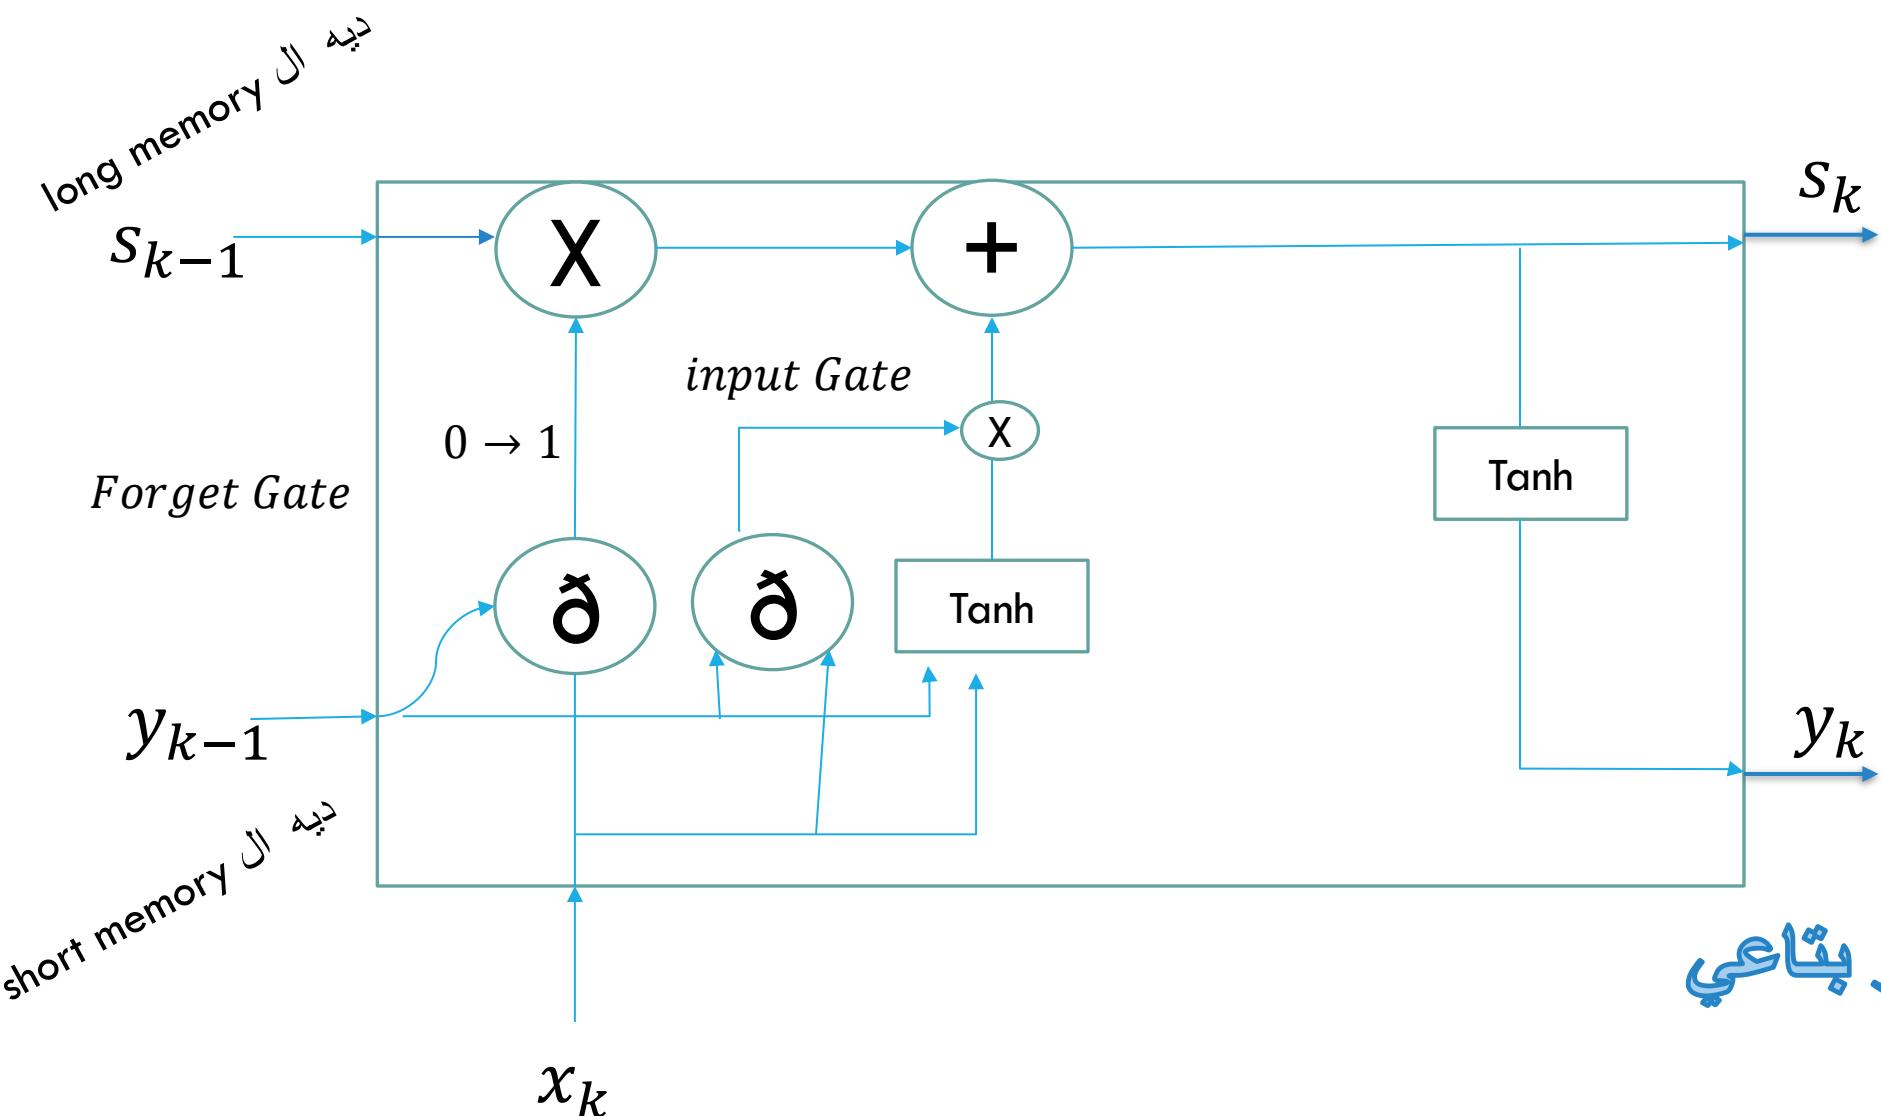


مین قال یعنی ان انت بناء على اللى انت شايفه والموقف الماضى روح ابديت الدفاتر
 فانا طيب جتنين فالقرار بتاع ال tanh اللى هيقولى اقطع علاقتك بيه قلبك هيقولك سامحوا
 فهاخذ الانبوت الحالى والانبوت السابق وادخله على sigmoid جديده اخلى الاوت بتاعها يضرب فى اوت التانش
 فديه هسميتها input gate اللى هو انبوت جديد للتاريخ بتاعى

- Cell state

Use outputs of previous gates to update current cell state

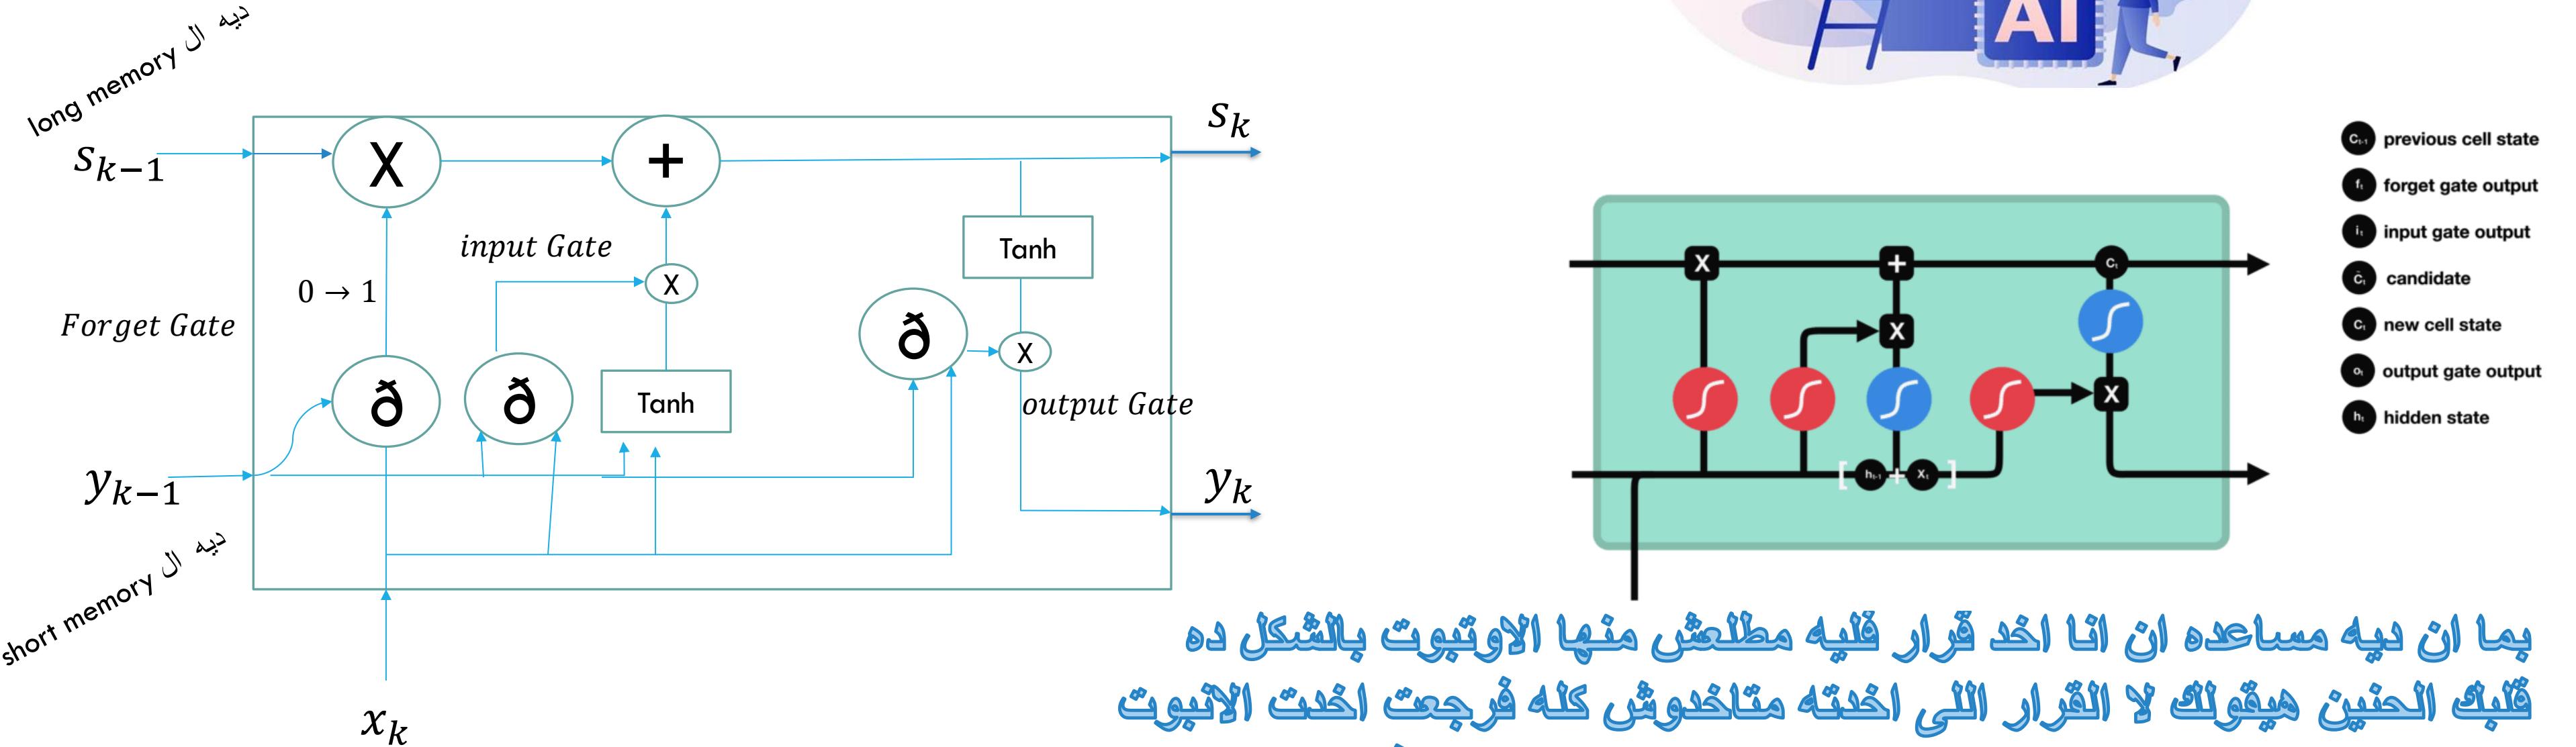
02 LSTM



بما ان دیه مساعده ان انا اخه قرار فلیه مطلعش منها الاوئیوت القرار بتاعی
پعنی ادخلوا على tanh بالشكل ده



02 LSTM

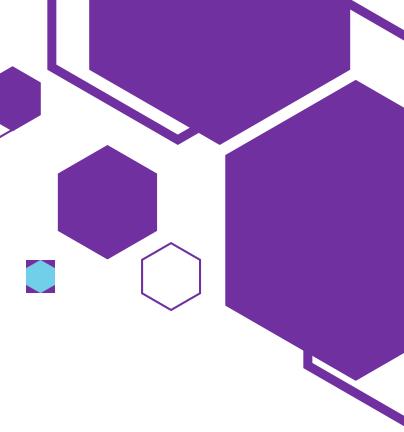
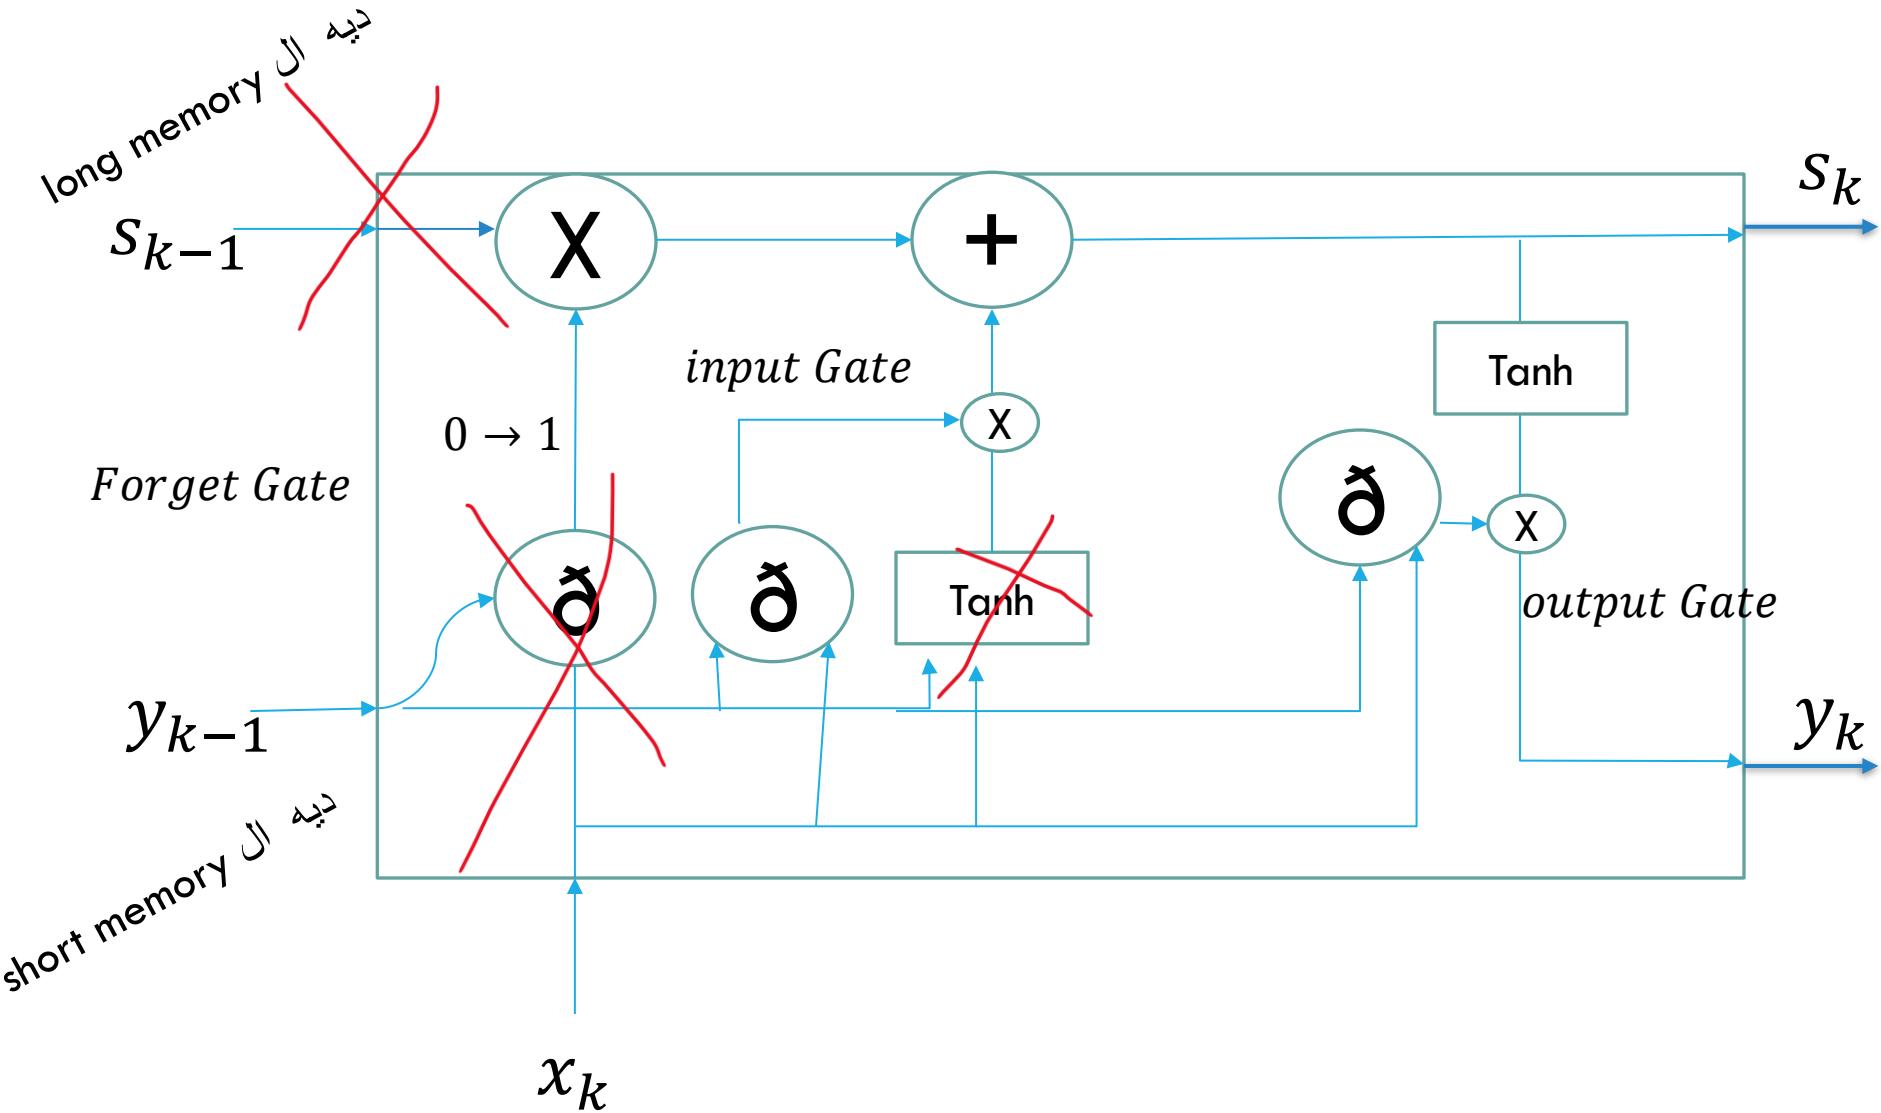


- c_{t-1} previous cell state
- f_t forget gate output
- i_t input gate output
- \tilde{c}_t candidate
- c_t new cell state
- o_t output gate output
- h_t hidden state

بما ان دیه مساعده ان انا اخذ قرار قلیه مطلعش منها الاوتبوت بالشكل ده
قلب الحنین هیقولك لا القرار اللي اخذته متاخدوش كله فرجعت اخذت الاوتبوت
مع الاوتبوت السابق دخلتهم على sigmoid وضربت الناتج في القرار بتاعي
ليه اللي هي ال output gate بتاعي.

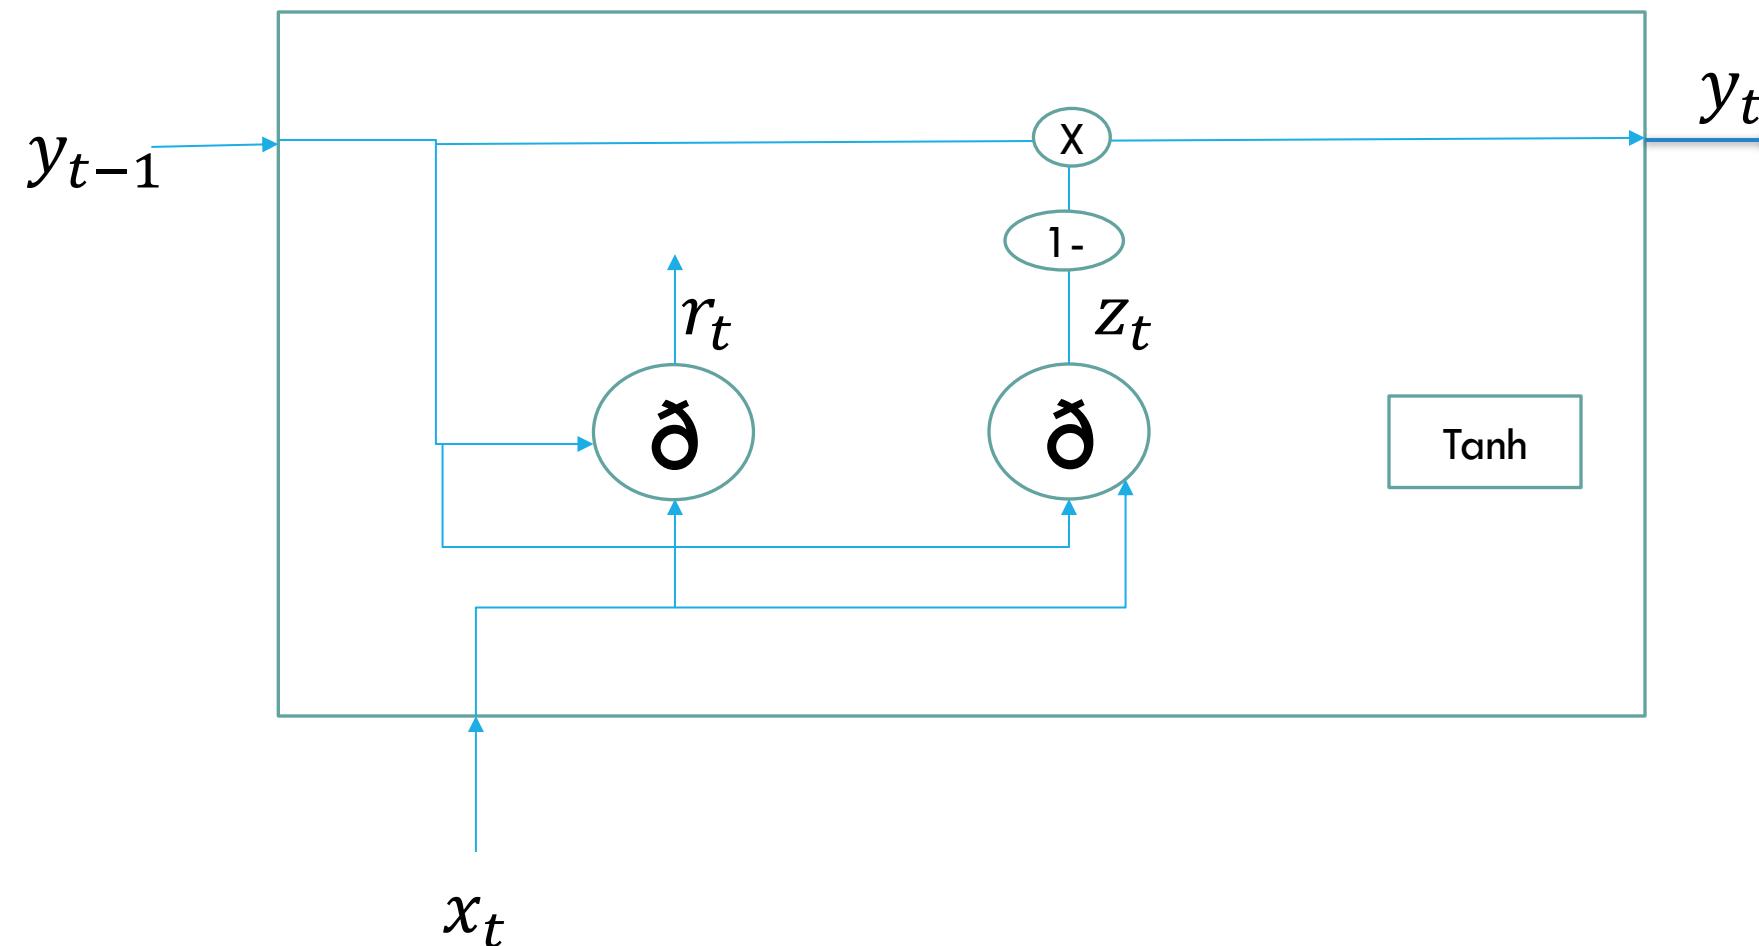
- Output gate
- Decides what new information to pass along as our hidden state

03 GRU



طب لیه التعقید ده کنه مش ممکن افول بس ان ($y_k = F(x_k, y_{k-1})$)
وکده کده ال y_{k-1} بتعتمد على ال y_{k-2} و هكذا....
بیکی هنلگی الجزء بتاع ال history وبالتبعیه هنلگی جزء ال sigmoid بتاعه
وکمان جزء ال tanh الی کان بیابدیت ال history

03 GRU

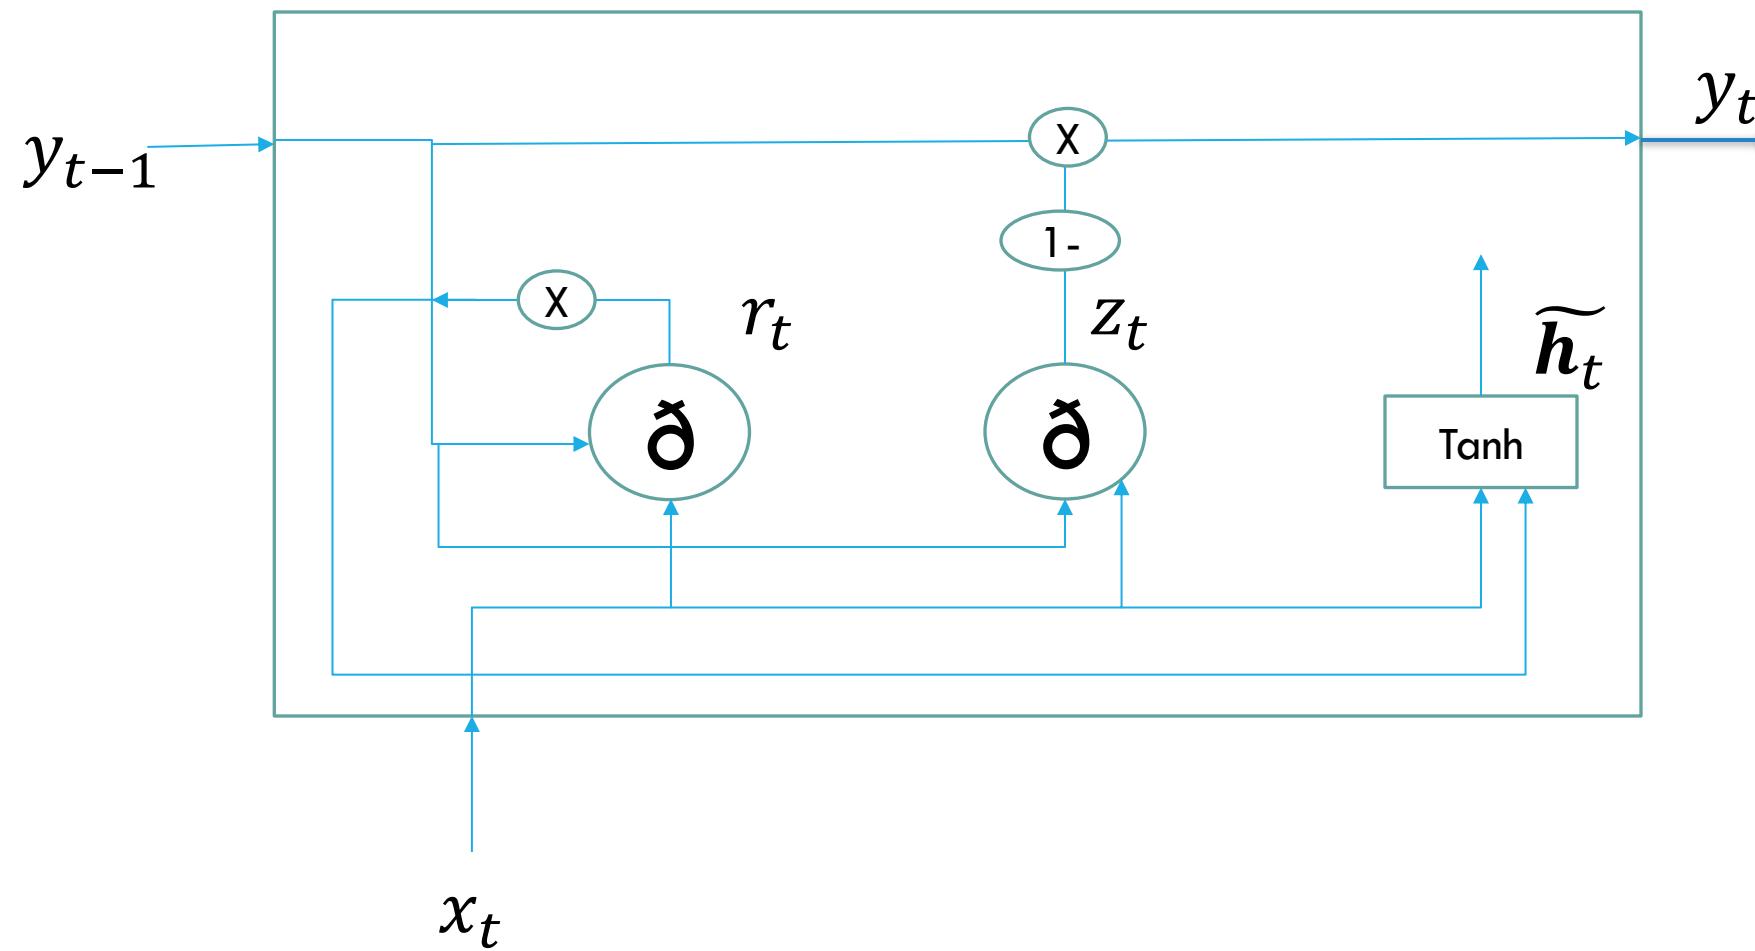


$$z_t = \sigma(y_{t-1}, x_t)$$

$$r_t = \sigma(y_{t-1}, x_t)$$

طب لیه التعویذ ده کله مش ممکن افول بس ان (y_{t-1}) وکده کده ال y_{t-1} بتعتمد على ال y_{t-2} ومهذا....
بیکی هنلگی الجزء بتاع ال history وبالتابعیه هنلگی جزء ال sigmoid بتاعه وکمان جزء ال tanh الی کان بیابدیت ال history

03 GRU



$$z_t = \sigma(y_{t-1}, x_t)$$

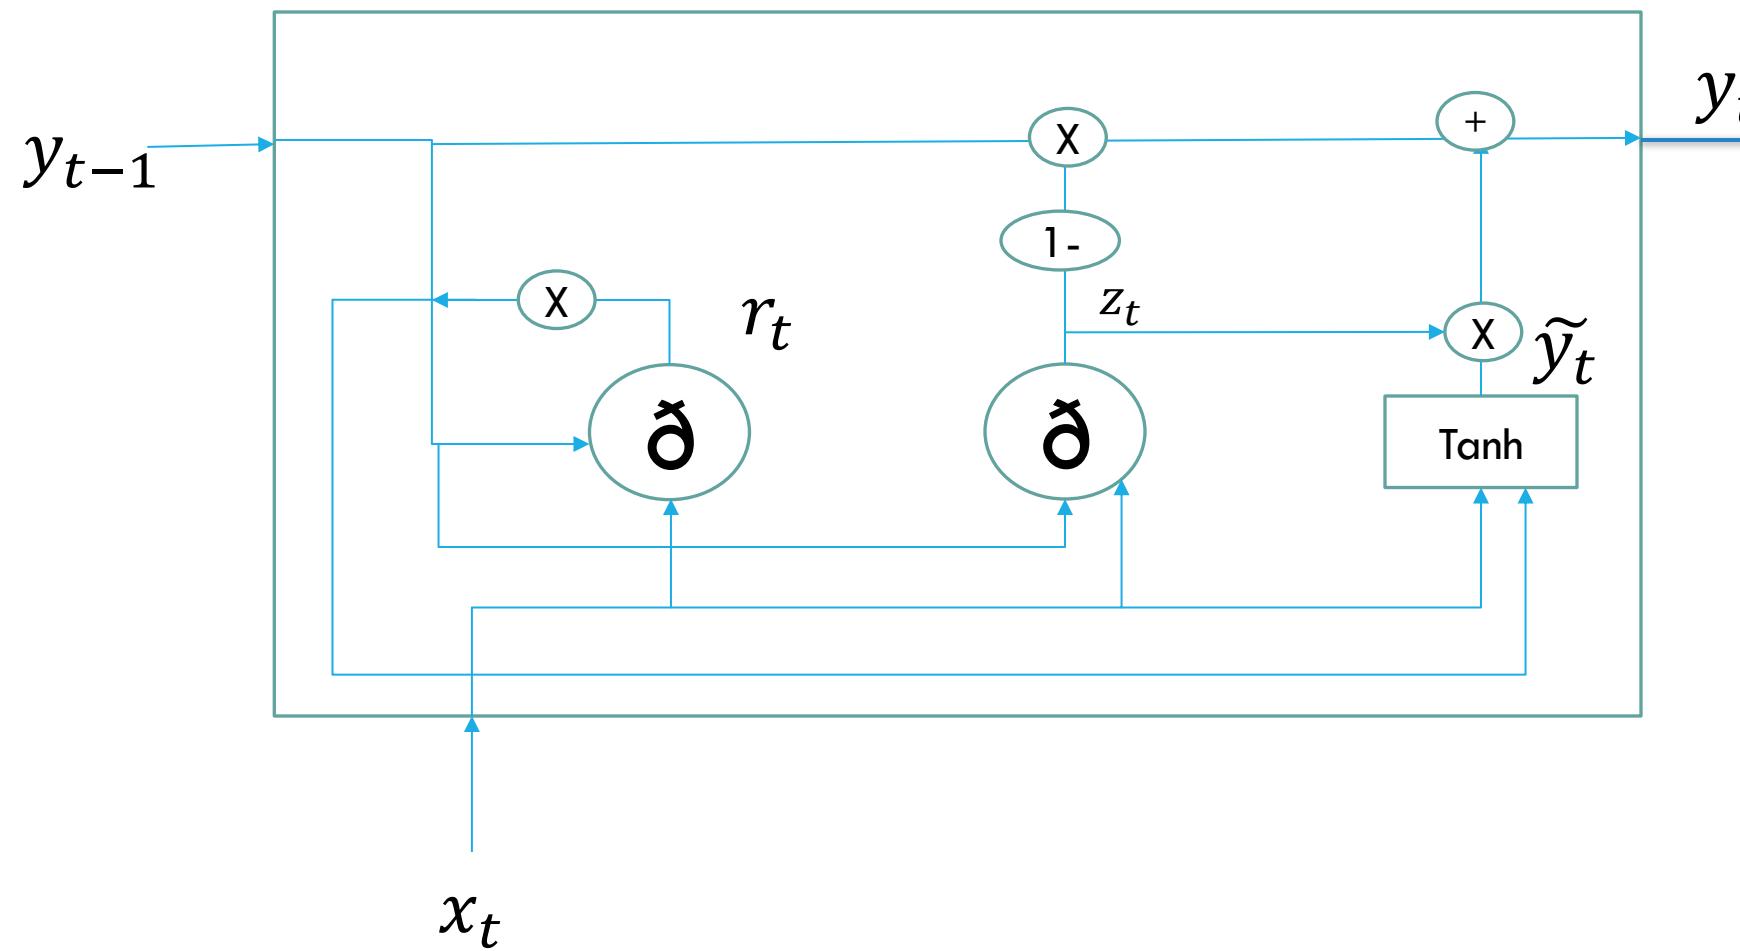
$$r_t = \sigma(y_{t-1}, x_t)$$

$$\tilde{h}_t = \tanh(r_t * y_{t-1}, x_t)$$

طب لیه التعویذ ده کله مش ممکن افول بس ان (y_{t-1}) وکده کده ال y_{t-1} بتعتمد على ال y_{t-2} ومهذا....
بیکی هنلگی الجزء بتاع ال history وبالتابعیه هنلگی جزء ال sigmoid بتاعه وکمان جزء ال \tanh الی کان بیابدیت ال history



03 GRU



$$z_t = \tilde{\sigma}(y_{t-1}, x_t)$$

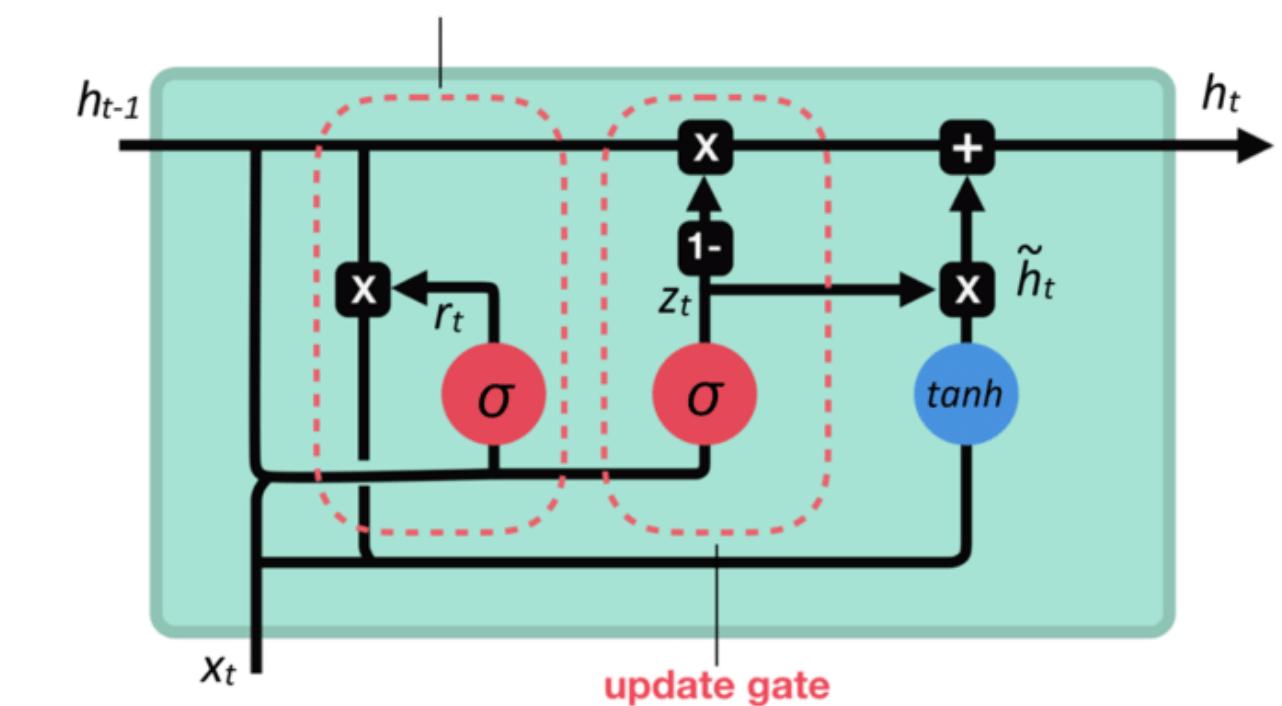
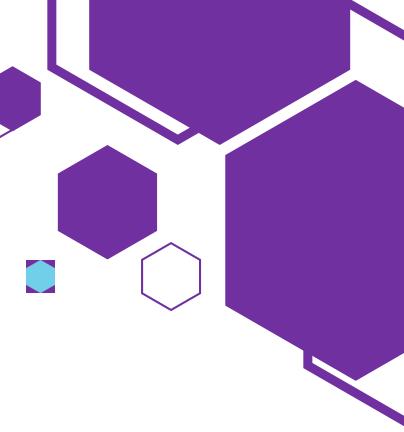
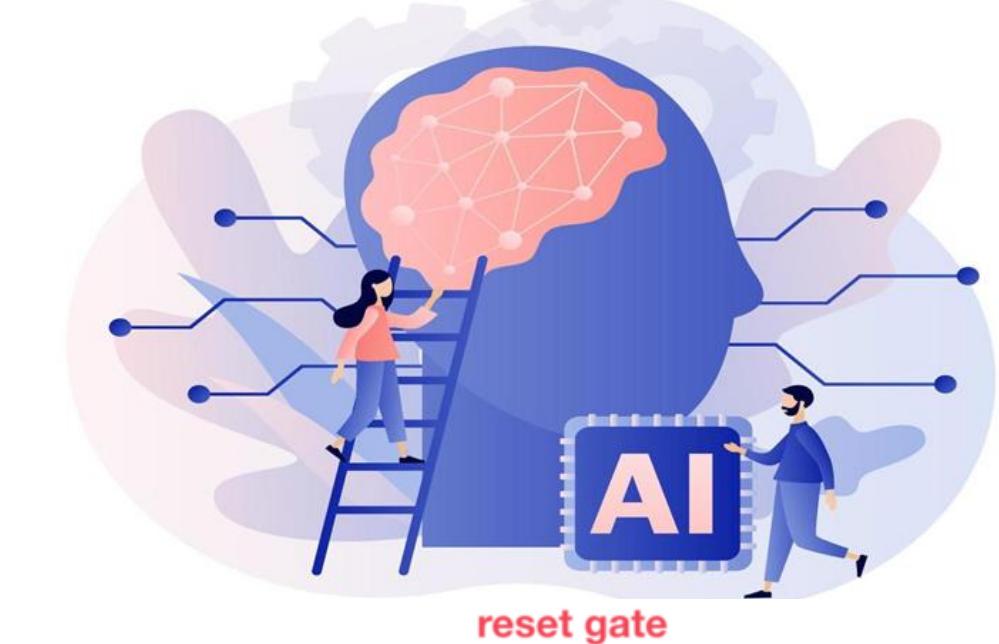
$$r_t = \tilde{\sigma}(y_{t-1}, x_t)$$

$$\tilde{y}_t = \text{Tanh}(r_t * y_{t-1} + (1 - z_t) * \tilde{y}_t)$$

$$y_t = (1 - z_t) * y_{t-1} + z_t * \tilde{y}_t$$

طب لیه التعقید ده کله مش ممکن افول بس ان ($y_t = F(x_t, y_{t-1})$) وکده کده ال y_{t-1} بتعتمد على ال y_{t-2} ومهذا....

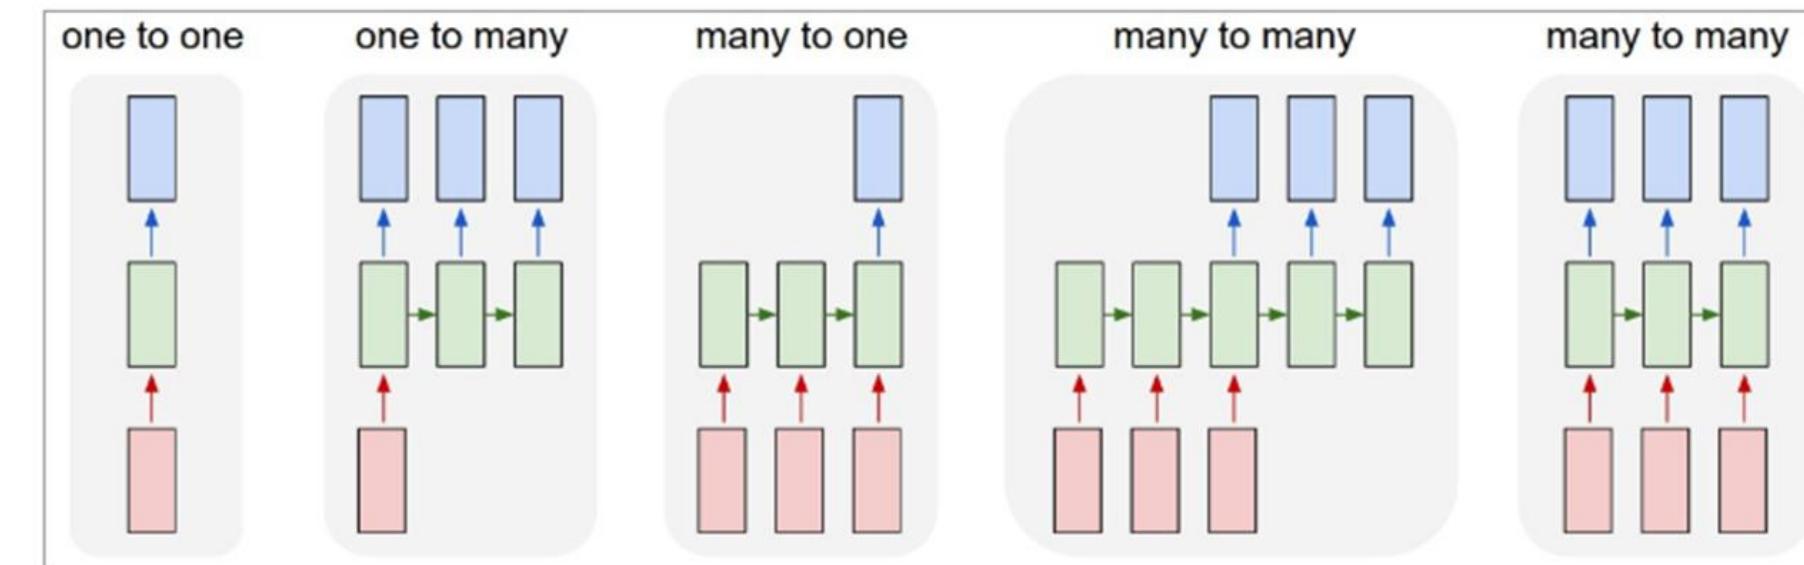
بیکی هنلگی الجزء بتاع ال history وبالاتبعیه هنلگی جزء ال sigmoid بتاعه وكمان جزء ال tanh الی کان بیابدیت ال history



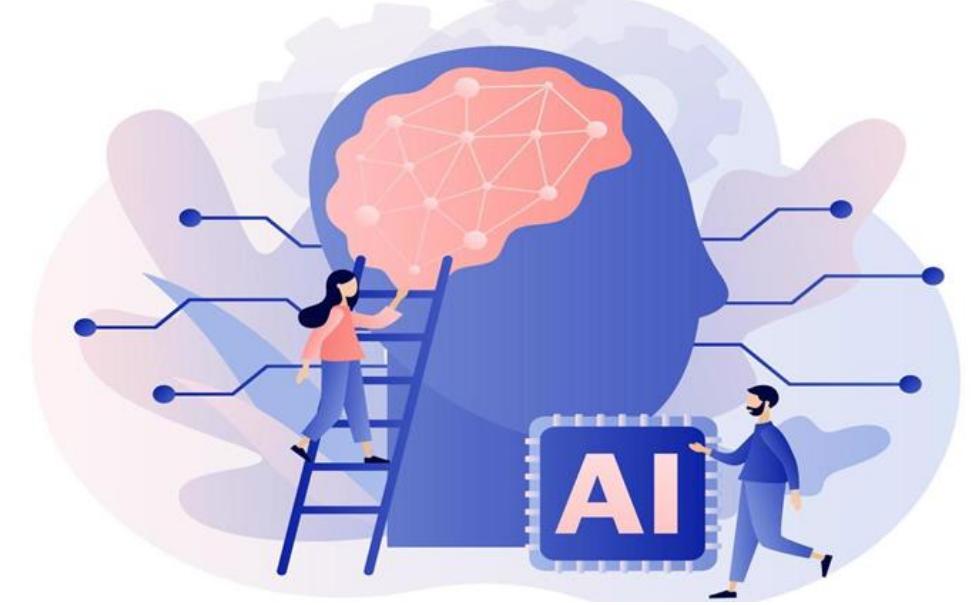
01 RNN topologies



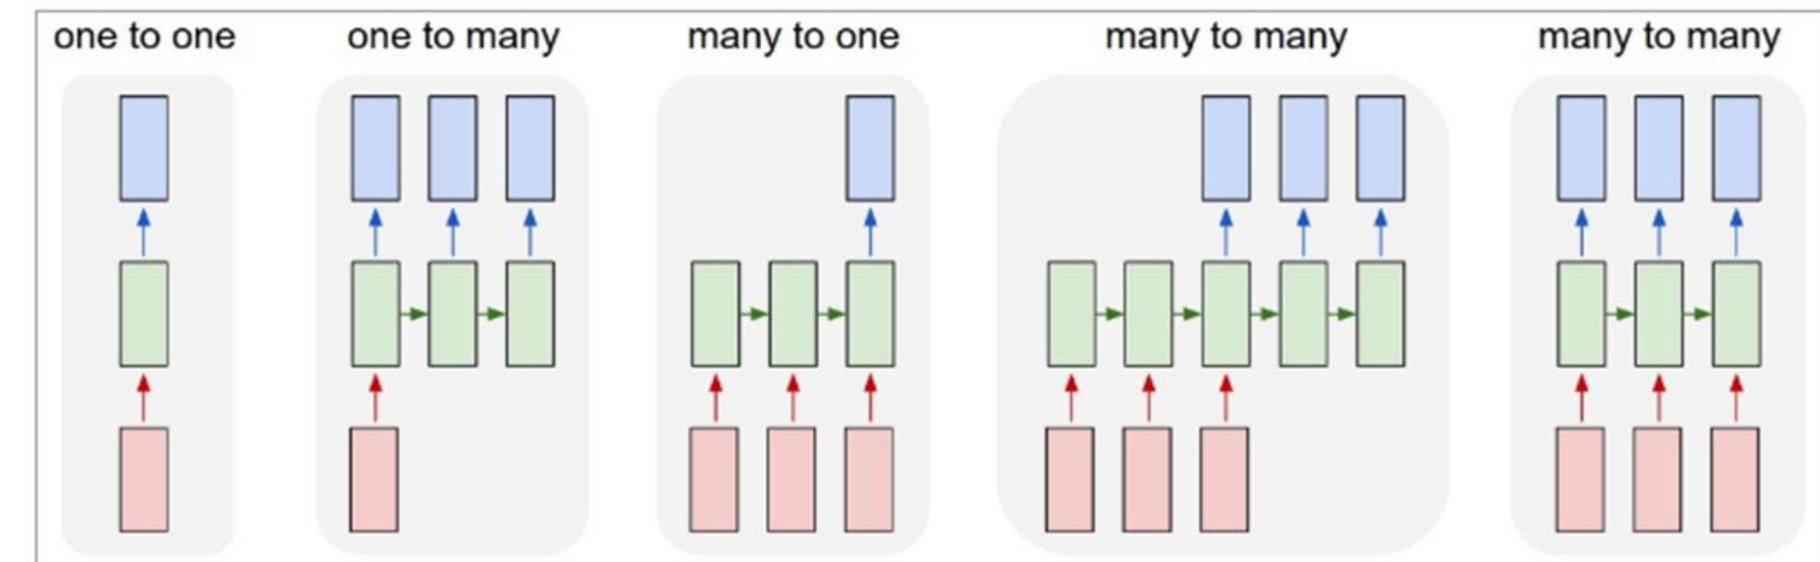
- The APIs for MLP and CNN architectures are limited. Both architectures accept a fixed-size tensor as input and produce a fixed-size tensor as output; and they perform the transformation from input to output in a fixed number of steps given by the number of layers in the model.
- RNNs don't have this limitation—you can have sequences in the input, the output, or both. This means that RNNs can be arranged in many ways to solve specific problems.
- RNNs combine the input vector with the previous state vector to produce a new state vector. This can be thought of as similar to running a program with some inputs and some internal variables.



01 RNN topologies



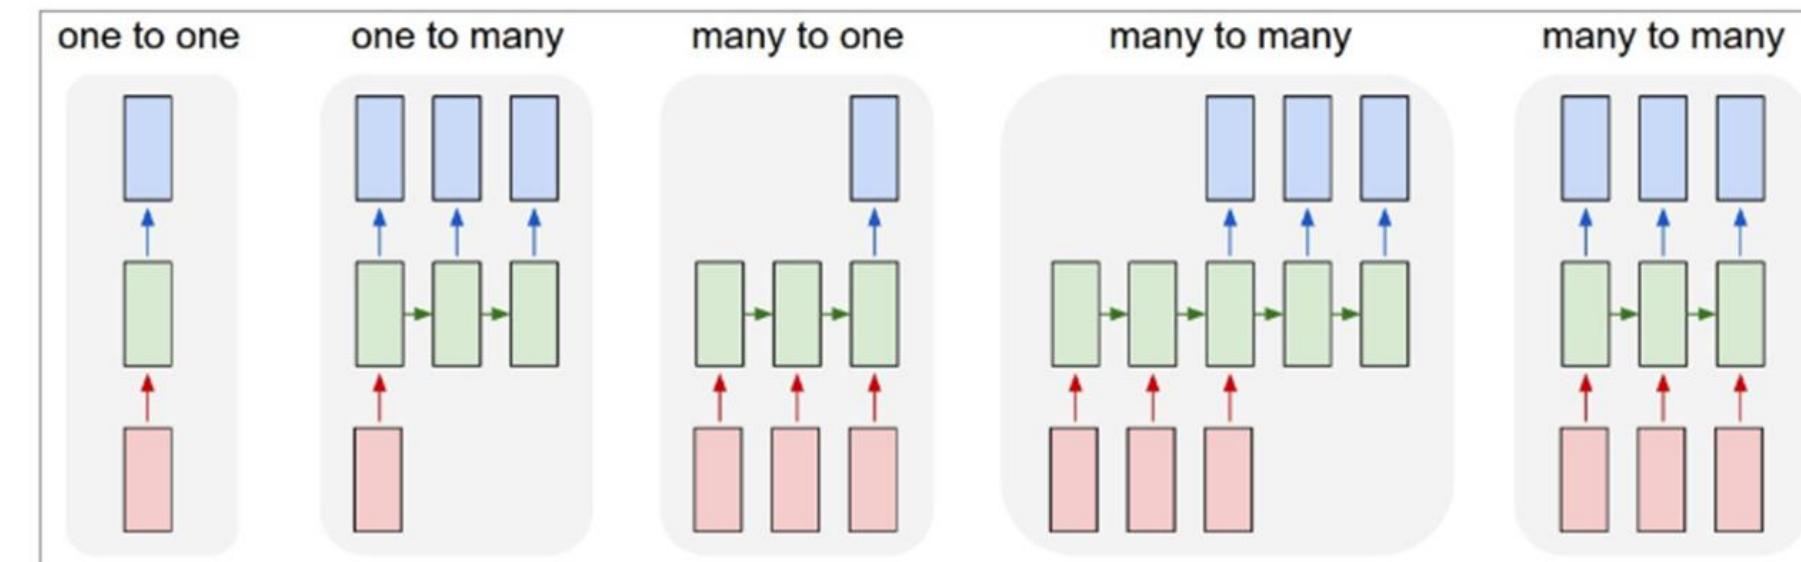
- The RNNs are more exciting because they allow us to operate over sequences of vectors: Sequences in the input, the output, or in the most general case both.
- This property of being able to work with sequences gives rise to a number of common topologies shown below:
- In the above diagram each rectangle is a vector and arrows represent functions (e.g. matrix multiply).
- Input vectors are in red, output vectors are in blue and green vectors hold the RNN's state.



01 RNN topologies



- From left to right:
 1. Vanilla mode of processing without RNN, from fixed-sized input to fixed-sized output (e.g. image classification).
 2. Sequence output (e.g. image captioning takes an image and outputs a sentence of words).
 3. Sequence input (e.g. sentiment analysis where a given sentence is classified as expressing positive or negative sentiment).
 4. Sequence input and sequence output (e.g. Machine Translation: an RNN reads a sentence in English and then outputs a sentence in French).
 5. Synced sequence input and output (e.g. video classification where we wish to label each frame of the video).



Thank You...!



End